

Nexus Family Pass - PostgreSQL Database Schema

Version: 1.0

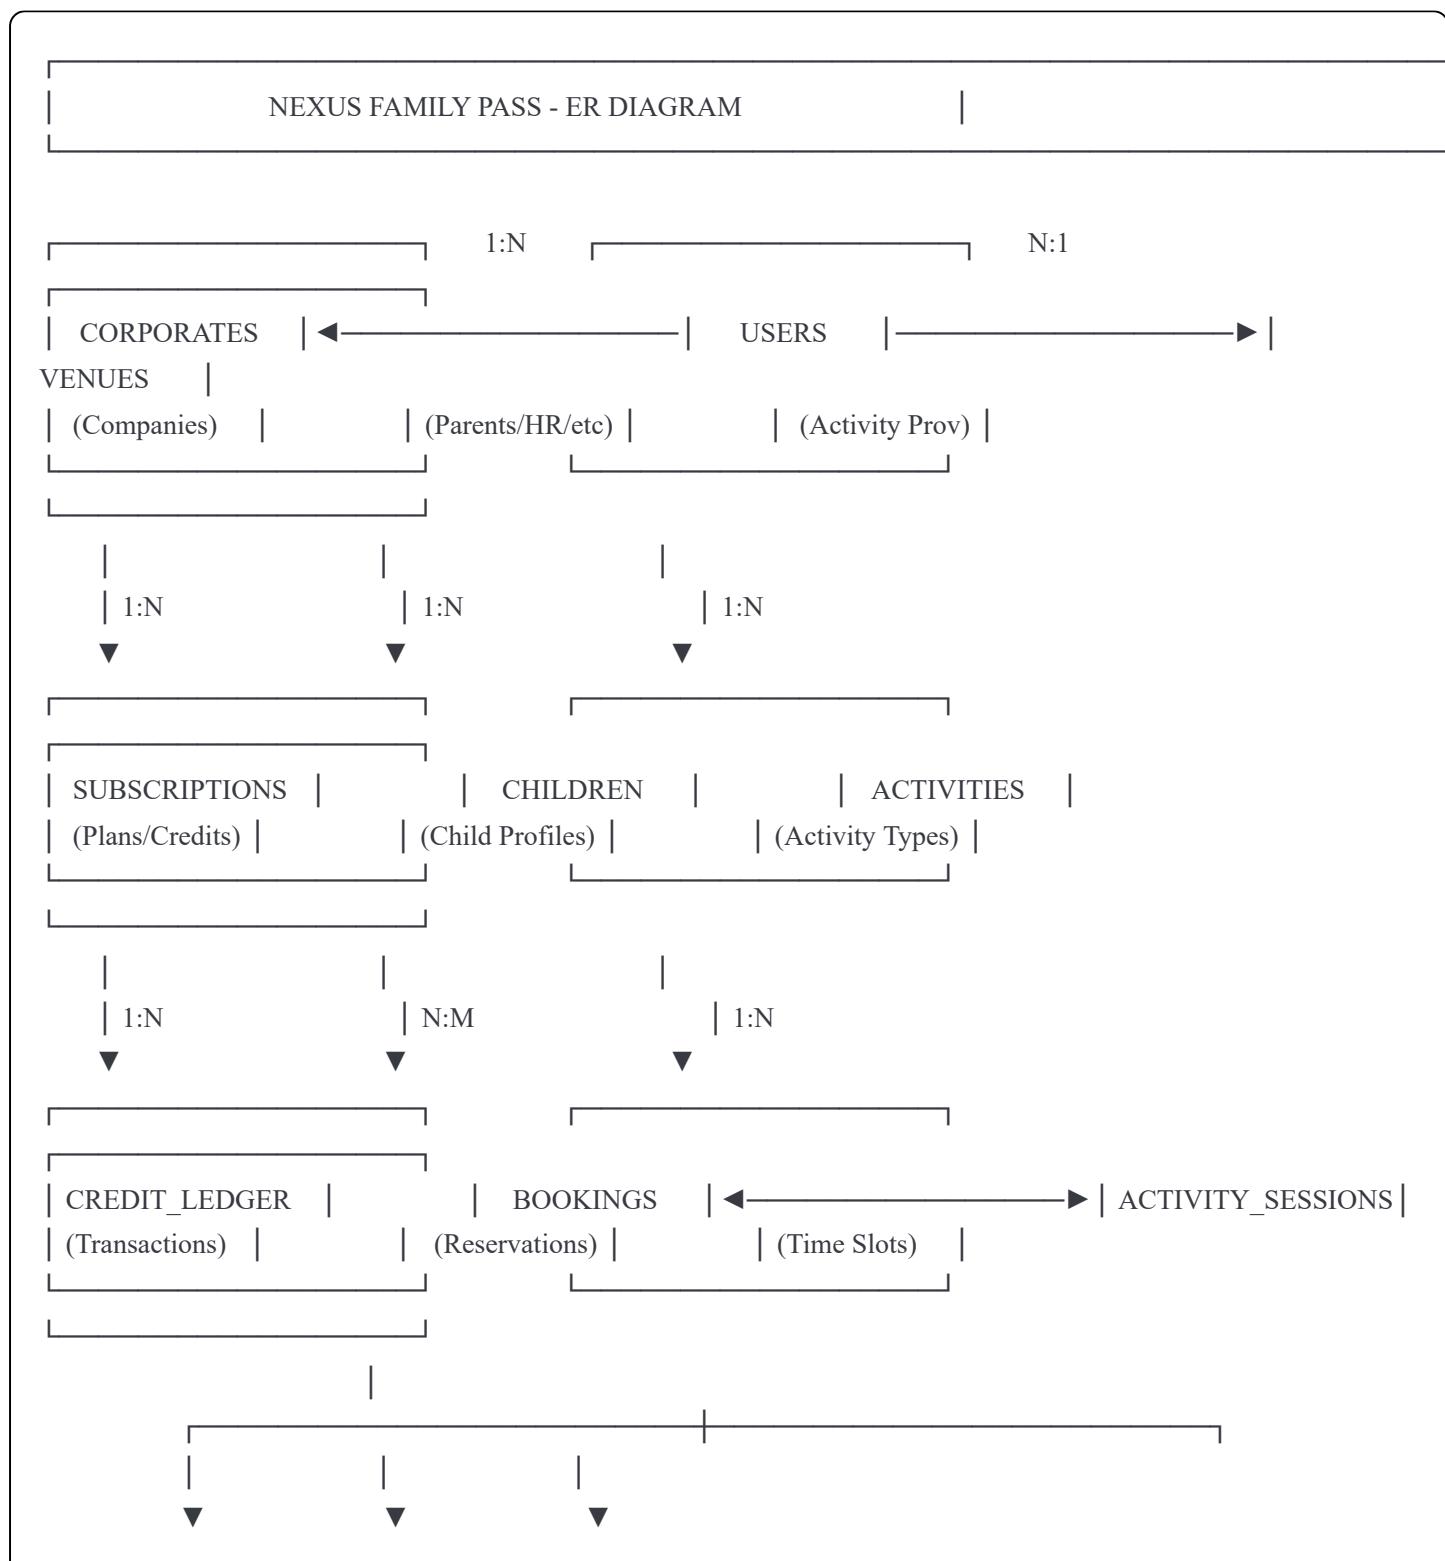
Last Updated: January 2026

Database: PostgreSQL 15+

Author: Senior Database Architect

1. Conceptual Entity-Relationship Diagram

Core Entities and Relationships





Additional Supporting Entities:



Relationship Summary

Relationship	Type	Description
Corporate → Users	One-to-Many	Each company has multiple employee users
Corporate → Subscriptions	One-to-Many	Companies can have multiple subscription periods
User (Parent) → Children	One-to-Many	Parents can register multiple children
Venue → Activities	One-to-Many	Venues offer multiple activity types
Activity → Sessions	One-to-Many	Activities have multiple scheduled sessions
Session → Bookings	One-to-Many	Sessions can have multiple child bookings
Child → Bookings	One-to-Many	Children can have multiple bookings
Booking → Feedback	One-to-One	Each booking can have one quick feedback
Booking → Review	One-to-One	Each booking can have one detailed review
User → Notifications	One-to-Many	Users receive multiple notifications
Venue → Venue Scores	One-to-Many	Weekly performance score history
Child → Waitlist	One-to-Many	Children can be on multiple waitlists
Child → AI Suggestions	One-to-Many	Monthly curated suggestions per child

1.5 Pinecone Integration Strategy

Division of Responsibilities

Data Type	PostgreSQL	Pinecone	Rationale
Activity metadata	<input checked="" type="checkbox"/> Store	<input type="checkbox"/> X	Relational queries, joins with venues
Activity embeddings	Reference ID only	<input checked="" type="checkbox"/> Store vectors	Semantic similarity search
Child preferences	Reference ID only	<input checked="" type="checkbox"/> Store vectors	Personalization matching
NL search queries	Audit log only	<input checked="" type="checkbox"/> Query vectors	Intent-based search
Sibling compatibility	Reference ID only	<input checked="" type="checkbox"/> Multi-vector ops	Group matching algorithms
Booking transactions	<input checked="" type="checkbox"/> Full storage	<input type="checkbox"/> X	ACID compliance required
Reviews/Ratings	<input checked="" type="checkbox"/> Full storage	<input type="checkbox"/> X	Aggregations, reporting
Venue scores	<input checked="" type="checkbox"/> Full storage	Metadata filter	Filtering during search

Pinecone Index Structure (Recommended)

```

Index: nexus-activities
├── Namespace: activity-embeddings
│   └── Vectors: Activity description embeddings
│       └── Metadata: {activity_id, venue_id, category, min_age, max_age,
│                         credits, is_messy, is_competitive, is_outdoor, venue_score}
|
├── Namespace: child-preferences
│   └── Vectors: Child interest/preference embeddings
│       └── Metadata: {child_id, parent_id, energy_level, social_pref, age}
|
└── Namespace: search-queries
    └── Vectors: Historical NL query embeddings (for improving search)
        └── Metadata: {query_id, user_id, resulted_in_booking, timestamp}

```

Sync Strategy

PostgreSQL remains the **source of truth** for all data. Pinecone vectors are derived and can be regenerated. The schema includes sync tracking columns to ensure consistency.

2. Detailed SQL Schema (DDL)

2.1 Extension Dependencies

```
sql  
-- =====  
-- NEXUS FAMILY PASS - DATABASE SCHEMA  
-- PostgreSQL 15+ Required  
-- =====  
  
-- Enable required PostgreSQL extensions for UUID generation and full-text search  
CREATE EXTENSION IF NOT EXISTS "pgcrypto";    -- Provides gen_random_uuid() function for generating UUIDs  
CREATE EXTENSION IF NOT EXISTS "pg_trgm";      -- Trigram extension for fuzzy text search capabilities  
CREATE EXTENSION IF NOT EXISTS "btree_gin";     -- GIN index support for efficient JSONB querying
```

2.2 Enum Types (Domain-Specific Value Constraints)

```
sql
```

```

-- =====
-- ENUM TYPE DEFINITIONS
-- Using PostgreSQL ENUM types ensures data integrity at the database level
-- and provides better query performance than VARCHAR with CHECK constraints
-- =====

-- User role enumeration defining the four primary user types in the system
-- Parent: End-user who books activities for their children
-- HR Admin: Corporate administrator managing employee benefits enrollment
-- Venue Admin: Activity provider managing their venue and listings
-- Platform Admin: System administrator with full platform access

CREATE TYPE user_role AS ENUM (
    'parent',          -- End-user parent who books activities for children
    'hr_admin',        -- HR administrator for corporate accounts
    'venue_admin',     -- Venue owner/manager who creates activities
    'platform_admin'   -- Nexus platform super administrator
);

-- Booking status tracking the complete lifecycle of a reservation
CREATE TYPE booking_status AS ENUM (
    'pending',         -- Booking created but awaiting venue confirmation
    'confirmed',       -- Booking confirmed by venue/system
    'completed',       -- Activity attended and marked complete
    'cancelled_parent', -- Cancelled by parent (refund rules apply)
    'cancelled_venue',  -- Cancelled by venue (full refund guaranteed)
    'no_show'          -- Child did not attend (credits forfeited)
);

-- Venue approval status for the venue onboarding workflow
CREATE TYPE venue_status AS ENUM (
    'pending_approval', -- Initial application submitted, awaiting review
    'active',           -- Approved and visible to parents
    'suspended',        -- Temporarily disabled due to issues
    'rejected'          -- Application denied (with feedback)
);

-- Credit transaction types for the financial ledger
CREATE TYPE credit_transaction_type AS ENUM (
    'allocation',      -- Monthly credits allocated from subscription
    'booking_debit',   -- Credits spent on booking an activity
    'refund_credit',   -- Credits returned due to cancellation (>48hrs)
    'partial_refund',  -- Partial credits returned when venue modifies activity (PRD 2.2)
    'forfeit',          -- Credits forfeited (cancellation <48hrs or no-show)
    'adjustment',       -- Manual adjustment by platform admin
    'expiry'            -- Monthly unused credits expired
);

```

-- Notification delivery channels supported by the platform

CREATE TYPE notification_channel **AS ENUM** (

- 'email', -- Primary email notification
- 'sms', -- SMS text message
- 'whatsapp', -- WhatsApp business messaging
- 'push', -- Mobile app push notification
- 'in_app' -- In-application notification center

);

-- Notification priority levels for delivery and display ordering

CREATE TYPE notification_priority **AS ENUM** (

- 'low', -- Informational, can be batched in digests
- 'normal', -- Standard delivery timing
- 'high', -- Time-sensitive, deliver immediately
- 'urgent' -- Critical alerts requiring immediate action

);

-- Child energy level preference from onboarding quiz (Question 1)

CREATE TYPE energy_level **AS ENUM** (

- 'calm.Focused', -- Prefers quiet, concentrated activities
- 'balanced', -- Moderate energy, flexible
- 'high.Energy' -- Active, physical activities preferred

);

-- Child social preference from onboarding quiz (Question 2)

CREATE TYPE social_preference **AS ENUM** (

- 'solo', -- Prefers individual activities
- 'small.Group', -- Comfortable in groups of 3-6
- 'big.Social' -- Thrives in larger group settings

);

-- Activity category classification for filtering and recommendations

CREATE TYPE activity_category **AS ENUM** (

- 'stem', -- Science, Technology, Engineering, Math
- 'arts.Crafts', -- Visual arts, crafting, creative projects
- 'music', -- Musical instruments, singing, composition
- 'sports', -- Physical sports and athletic activities
- 'nature', -- Outdoor exploration, environmental education
- 'cooking', -- Culinary skills, baking, food preparation
- 'reading', -- Literature, storytelling, book clubs
- 'building', -- Construction, LEGO, architecture
- 'dance', -- Dance styles, movement, choreography
- 'drama' -- Theater, acting, performance arts

);

-- Subscription plan tiers for corporate accounts

```
CREATE TYPE subscription_tier AS ENUM (
    'basic',      -- Entry-level plan with limited credits
    'standard',   -- Mid-tier with moderate credits
    'premium',    -- High-tier with premium venue access
    'enterprise' -- Custom enterprise agreements
);
```

-- Cancellation reason categories (from UI Spec line 330)

```
CREATE TYPE cancellation_reason_type AS ENUM (
    'schedule_conflict', -- Parent has scheduling conflict
    'child_sick',        -- Child is unwell
    'found_better_option', -- Found a preferable activity
    'venue_cancelled',   -- Venue initiated cancellation
    'weather',           -- Weather-related cancellation
    'transportation',    -- Transportation issues
    'other'              -- Other reason (with notes)
);
```

-- Time of day for session filtering (from UI Spec line 196)

```
CREATE TYPE time_of_day AS ENUM (
    'morning',     -- Before 12:00 PM
    'afternoon',   -- 12:00 PM - 5:00 PM
    'evening'      -- After 5:00 PM
);
```

-- Employee invitation status tracking

```
CREATE TYPE invitation_status AS ENUM (
    'pending',      -- Invitation sent, awaiting response
    'accepted',    -- Employee accepted and registered
    'expired',     -- Invitation link expired
    'revoked'      -- HR revoked the invitation
);
```

2.3 Core User and Authentication Tables

sql

```

-- =====
-- CORPORATE ACCOUNTS TABLE
-- Stores B2B company information for the corporate subscription model
-- Each company purchases activity credits for their employees
-- =====

CREATE TABLE corporates (
    -- Primary identifier using UUID for security and distributed systems compatibility
    id UUID PRIMARY KEY DEFAULT gen_random_uuid(),

    -- Company legal name as registered for invoicing purposes
    company_name VARCHAR(255) NOT NULL,

    -- Trading name or brand name if different from legal name
    display_name VARCHAR(255),

    -- Corporate tax identification number for B2B invoicing compliance
    tax_id VARCHAR(50),

    -- Primary business industry for analytics and reporting
    industry VARCHAR(100),

    -- Total number of employees at the company (for plan sizing)
    employee_count INTEGER CHECK (employee_count > 0),

    -- Corporate headquarters address for contract purposes
    billing_address_line1 VARCHAR(255),
    billing_address_line2 VARCHAR(255),
    billing_city VARCHAR(100),
    billing_state VARCHAR(100),
    billing_postal_code VARCHAR(20),
    billing_country VARCHAR(100) DEFAULT 'USA',

    -- Primary contact person for the corporate account
    primary_contact_name VARCHAR(255) NOT NULL,
    primary_contact_email VARCHAR(255) NOT NULL,
    primary_contact_phone VARCHAR(50),

    -- Stripe Connect customer ID for payment processing
    stripe_customer_id VARCHAR(255),

    -- Payment terms in days (default Net-30 as per PRD)
    payment_terms_days INTEGER DEFAULT 30 CHECK (payment_terms_days >= 0),

    -- Enterprise procurement system integration identifiers
    sap_vendor_id VARCHAR(100),      -- SAP integration ID
    coupa_supplier_id VARCHAR(100),   -- Coupa integration ID
);

```

```

-- Account status tracking
is_active BOOLEAN DEFAULT true,

-- Contract dates for subscription management
contract_start_date DATE NOT NULL,
contract_end_date DATE,

-- Audit timestamps
created_at TIMESTAMP WITH TIME ZONE DEFAULT CURRENT_TIMESTAMP,
updated_at TIMESTAMP WITH TIME ZONE DEFAULT CURRENT_TIMESTAMP
);

-- Index on company name for search and autocomplete functionality
CREATE INDEX idx_corporates_company_name ON corporates(company_name);

-- Index on active status for filtering active corporate accounts
CREATE INDEX idx_corporates_is_active ON corporates(is_active);

-- Index on contract dates for renewal tracking queries
CREATE INDEX idx_corporates_contract_dates ON corporates(contract_start_date, contract_end_date);

-- Add table comment for documentation
COMMENT ON TABLE corporates IS 'B2B corporate accounts that purchase activity subscriptions for employees';

```

```

-- =====
-- USERS TABLE
-- Unified user table supporting all four user roles: Parent, HR Admin,
-- Venue Admin, and Platform Admin with role-based access control
-- =====

CREATE TABLE users (
    -- Primary identifier using UUID for security
    id UUID PRIMARY KEY DEFAULT gen_random_uuid(),

    -- User email address - used for login and notifications (must be unique)
    email VARCHAR(255) NOT NULL UNIQUE,

    -- Argon2/bcrypt hashed password - NULL for SSO-only users
    password_hash VARCHAR(255),

    -- User's role determining their access level and UI experience
    role user_role NOT NULL,

    -- User's display name shown in the UI
    first_name VARCHAR(100) NOT NULL,
    last_name VARCHAR(100) NOT NULL,

    -- User's phone number for SMS/WhatsApp notifications

```

```
phone VARCHAR(50),  
  
-- URL or path to user's profile avatar image  
avatar_url VARCHAR(500),  
  
-- Foreign key to corporate account (required for parents and HR admins)  
corporate_id UUID REFERENCES corporates(id) ON DELETE SET NULL,  
  
-- Foreign key to venue (required for venue admins, defined after venues table)  
-- venue_id UUID REFERENCES venues(id) ON DELETE SET NULL,  
  
-- Department within the corporate for HR reporting aggregation  
department VARCHAR(100),  
  
-- SSO configuration for corporate users  
sso_provider VARCHAR(50),          -- 'saml' or 'oidc'  
sso_subject_id VARCHAR(255),       -- Unique ID from SSO provider  
  
-- Multi-factor authentication settings  
mfa_enabled BOOLEAN DEFAULT false,  
mfa_secret VARCHAR(255),           -- TOTP secret (encrypted at app level)  
  
-- Account status flags  
is_active BOOLEAN DEFAULT true,  
is_email_verified BOOLEAN DEFAULT false,  
email_verified_at TIMESTAMP WITH TIME ZONE,  
  
-- Password reset token management  
password_reset_token VARCHAR(255),  
password_reset_expires_at TIMESTAMP WITH TIME ZONE,  
  
-- Last login tracking for security and engagement metrics  
last_login_at TIMESTAMP WITH TIME ZONE,  
last_login_ip INET,  
  
-- Notification preferences stored as flexible JSONB for extensibility  
-- Structure: { "email": true, "sms": false, "push": true, "frequency": "immediate" }  
notification_preferences JSONB DEFAULT '{  
    "email": true,  
    "sms": false,  
    "whatsapp": false,  
    "push": true,  
    "frequency": "immediate"  
}':jsonb,  
  
-- User's timezone for displaying local times  
timezone VARCHAR(50) DEFAULT 'America/New_York',
```

```

-- Preferred language for UI and notifications
preferred_language VARCHAR(10) DEFAULT 'en',

-- Remember me token for persistent login sessions (UI Spec line 42)
remember_token VARCHAR(255),
remember_token_expires_at TIMESTAMP WITH TIME ZONE,

-- Coworker matching consent (PRD Feature 4 & UI Spec line 413)
-- Allows children to be matched with coworkers' children for group activities
allow_coworker_matching BOOLEAN DEFAULT false,
coworker_matching_consent_date TIMESTAMP WITH TIME ZONE,

-- Audit timestamps
created_at TIMESTAMP WITH TIME ZONE DEFAULT CURRENT_TIMESTAMP,
updated_at TIMESTAMP WITH TIME ZONE DEFAULT CURRENT_TIMESTAMP
);

-- Index on email for login queries (already has unique constraint)
-- Index on role for filtering users by type
CREATE INDEX idx_users_role ON users(role);

-- Index on corporate_id for finding all users in a company
CREATE INDEX idx_users_corporate_id ON users(corporate_id);

-- Composite index for active users by role (common dashboard query)
CREATE INDEX idx_users_active_role ON users(is_active, role);

-- Index on SSO identifiers for SSO login lookup
CREATE INDEX idx_users_sso ON users(sso_provider, sso_subject_id) WHERE sso_provider IS NOT NULL;
-- GIN index on notification_preferences for JSONB queries
CREATE INDEX idx_users_notification_prefs ON users USING GIN (notification_preferences);

COMMENT ON TABLE users IS 'Unified user accounts supporting parents, HR admins, venue admins, and platform admins';

-----  

-- PARENT LOCATION PREFERENCES TABLE
-- Stores multiple addresses (home, work, school) for geo-fence filtering
-- when recommending activities within 15-minute drive distance
-----  

CREATE TABLE user_locations (
    -- Primary identifier
    id UUID PRIMARY KEY DEFAULT gen_random_uuid(),
    -- Reference to the parent user (only parents need location preferences)
    user_id UUID NOT NULL REFERENCES users(id) ON DELETE CASCADE,
    -- Location type identifier
    location_type VARCHAR(20) NOT NULL CHECK (location_type IN ('home', 'work', 'school', 'other')),

```

```

-- Human-readable label for the location
label VARCHAR(100),

-- Full address components for display
address_line1 VARCHAR(255) NOT NULL,
address_line2 VARCHAR(255),
city VARCHAR(100) NOT NULL,
state VARCHAR(100),
postal_code VARCHAR(20) NOT NULL,
country VARCHAR(100) DEFAULT 'USA',

-- Geocoded coordinates for distance calculations (stored as separate fields)
-- Note: For production, consider PostGIS GEOGRAPHY type for proper spatial queries
latitude DECIMAL(10, 8),          -- Range: -90 to +90
longitude DECIMAL(11, 8),         -- Range: -180 to +180

-- Flag indicating if this is the primary address for search radius
is_primary BOOLEAN DEFAULT false,

-- Preferred search radius in minutes driving time (default 15 per PRD)
default_search_radius_minutes INTEGER DEFAULT 15 CHECK (default_search_radius_minutes BETWEEN 5 AND 60)

-- Audit timestamps
created_at TIMESTAMP WITH TIME ZONE DEFAULT CURRENT_TIMESTAMP,
updated_at TIMESTAMP WITH TIME ZONE DEFAULT CURRENT_TIMESTAMP,

-- Ensure only one location of each type per user
CONSTRAINT unique_user_location_type UNIQUE (user_id, location_type)
);

-- Index on user_id for fetching all locations for a user
CREATE INDEX idx_user_locations_user_id ON user_locations(user_id);
-- Index on coordinates for proximity queries
CREATE INDEX idx_user_locations_coords ON user_locations(latitude, longitude) WHERE latitude IS NOT NULL;

COMMENT ON TABLE user_locations IS 'Parent address preferences for geo-fence filtering of activity recommendations';

```

2.4 Venue and Activity Tables

sql

```
-- =====
-- VENUES TABLE
-- Activity providers (gymnasiums, art studios, STEM centers, etc.)
-- Subject to vetting and approval process before going live
-- =====

CREATE TABLE venues (
    -- Primary identifier
    id UUID PRIMARY KEY DEFAULT gen_random_uuid(),
    -- Venue's business name as displayed to parents
    name VARCHAR(255) NOT NULL,
    -- URL-friendly slug for venue profile pages
    slug VARCHAR(255) NOT NULL UNIQUE,
    -- Brief tagline for venue cards (150 char limit per UI spec)
    short_description VARCHAR(150),
    -- Full venue description with amenities and philosophy
    full_description TEXT,
    -- Primary activity category this venue specializes in
    primary_category activity_category,
    -- Venue approval and visibility status
    status venue_status DEFAULT 'pending_approval',
    -- Physical address of the venue
    address_line1 VARCHAR(255) NOT NULL,
    address_line2 VARCHAR(255),
    city VARCHAR(100) NOT NULL,
    state VARCHAR(100) NOT NULL,
    postal_code VARCHAR(20) NOT NULL,
    country VARCHAR(100) DEFAULT 'USA',
    -- Geocoded coordinates for distance calculations
    latitude DECIMAL(10, 8) NOT NULL,
    longitude DECIMAL(11, 8) NOT NULL,
    -- Contact information
    contact_email VARCHAR(255) NOT NULL,
    contact_phone VARCHAR(50) NOT NULL,
    website_url VARCHAR(500),
    -- Business hours stored as JSONB for flexibility
    -- Structure: { "monday": {"open": "09:00", "close": "18:00"}, ... }
```

```
business_hours JSONB DEFAULT '{}':jsonb,  
  
-- Media assets  
logo_url VARCHAR(500),  
-- Array of image URLs for venue gallery  
gallery_images JSONB DEFAULT '[]':jsonb,  
  
-- Insurance compliance (required per PRD Section 5.3)  
insurance_provider VARCHAR(255),  
insurance_policy_number VARCHAR(100),  
insurance_coverage_amount DECIMAL(12, 2) CHECK (insurance_coverage_amount >= 1000000),  
insurance_expiry_date DATE,  
insurance_document_url VARCHAR(500),  
  
-- Background check compliance tracking  
background_check_completed BOOLEAN DEFAULT false,  
background_check_date DATE,  
  
-- Safety audit compliance  
safety_audit_completed BOOLEAN DEFAULT false,  
safety_audit_date DATE,  
safety_audit_score INTEGER CHECK (safety_audit_score BETWEEN 0 AND 100),  
  
-- Accessibility features available at the venue  
accessibility_features JSONB DEFAULT '[]':jsonb,  
  
-- Banking information for payouts (encrypted at app level)  
payout_bank_name VARCHAR(255),  
payout_account_last_four VARCHAR(4),  
stripe_connect_account_id VARCHAR(255),  
  
-- Platform fee percentage for this venue (default 15-20% per PRD)  
platform_fee_percentage DECIMAL(5, 2) DEFAULT 17.50 CHECK (platform_fee_percentage BETWEEN 0 AND 100),  
  
-- Block booking configuration  
-- Number of days in advance the platform pre-books slots  
block_booking_advance_days INTEGER DEFAULT 30,  
-- Hours before activity when unsold inventory is released back  
inventory_release_hours INTEGER DEFAULT 72,  
  
-- Current computed performance score (0-100, recalculated weekly)  
current_performance_score DECIMAL(5, 2) DEFAULT 50.00 CHECK (current_performance_score BETWEEN 0 AND 100),  
  
-- Flag indicating if venue is approved for proactive curation (score ≥60)  
is_curation_eligible BOOLEAN DEFAULT false,  
  
-- Annual re-verification tracking
```

```

last_verification_date DATE,
next_verification_due DATE,

-- Admin user who manages this venue
primary_admin_user_id UUID,

-- Rejection feedback if application was denied
rejection_reason TEXT,
rejection_date TIMESTAMP WITH TIME ZONE,

-- Audit timestamps
created_at TIMESTAMP WITH TIME ZONE DEFAULT CURRENT_TIMESTAMP,
updated_at TIMESTAMP WITH TIME ZONE DEFAULT CURRENT_TIMESTAMP,
approved_at TIMESTAMP WITH TIME ZONE
);

-- Index on slug for URL lookups
CREATE INDEX idx_venues_slug ON venues(slug);

-- Index on status for filtering pending/active venues
CREATE INDEX idx_venues_status ON venues(status);

-- Spatial index on coordinates for proximity searches
CREATE INDEX idx_venues_coordinates ON venues(latitude, longitude);

-- Index on performance score for curation filtering (exclude <60)
CREATE INDEX idx_venues_performance ON venues(current_performance_score) WHERE status = 'active';

-- Index on insurance expiry for compliance alerts
CREATE INDEX idx_venues_insurance_expiry ON venues(insurance_expiry_date);

-- Index on category for filtering
CREATE INDEX idx_venues_category ON venues(primary_category);

-- GIN index on accessibility features for JSONB array queries
CREATE INDEX idx_venues_accessibility ON venues USING GIN (accessibility_features);

COMMENT ON TABLE venues IS 'Activity provider venues subject to vetting, approval, and performance scoring';

-- Now add the venue_id foreign key to users table
ALTER TABLE users ADD COLUMN venue_id UUID REFERENCES venues(id) ON DELETE SET NULL;
CREATE INDEX idx_users_venue_id ON users(venue_id) WHERE venue_id IS NOT NULL;

=====

-- VENUE PERFORMANCE SCORES HISTORY TABLE
-- Weekly snapshots of venue performance metrics for trend analysis
-- Score components weighted per PRD: Parent feedback 40%, Repeat booking 25%,
-- Cancellation rate 20% (negative), No-show rate 15%
=====

CREATE TABLE venue_performance_scores (
    -- Primary identifier

```

```

id UUID PRIMARY KEY DEFAULT gen_random_uuid(),

-- Reference to the venue being scored
venue_id UUID NOT NULL REFERENCES venues(id) ON DELETE CASCADE,

-- Week start date for this score snapshot (always Monday)
score_week_start DATE NOT NULL,

-- Individual component scores (0-100 scale)
parent_feedback_score DECIMAL(5, 2) CHECK (parent_feedback_score BETWEEN 0 AND 100),
repeat_booking_score DECIMAL(5, 2) CHECK (repeat_booking_score BETWEEN 0 AND 100),
cancellation_score DECIMAL(5, 2) CHECK (cancellation_score BETWEEN 0 AND 100),
no_show_score DECIMAL(5, 2) CHECK (no_show_score BETWEEN 0 AND 100),

-- Raw metrics used in calculation (for transparency and debugging)
total_bookings INTEGER DEFAULT 0,
total_reviews INTEGER DEFAULT 0,
average_rating DECIMAL(3, 2),
repeat_booking_rate DECIMAL(5, 2),
venue_cancellation_rate DECIMAL(5, 2),
parent_no_show_rate DECIMAL(5, 2),

-- Final weighted composite score
composite_score DECIMAL(5, 2) NOT NULL CHECK (composite_score BETWEEN 0 AND 100),

-- Score change from previous week
score_change DECIMAL(5, 2),

-- AI-generated improvement suggestions stored as JSONB array
improvementSuggestions JSONB DEFAULT '[]'::jsonb,

-- Timestamp when this score was calculated
calculated_at TIMESTAMP WITH TIME ZONE DEFAULT CURRENT_TIMESTAMP,

-- Ensure one score per venue per week
CONSTRAINT unique_venue_weekly_score UNIQUE (venue_id, score_week_start)
);

-- Index for retrieving score history for a venue
CREATE INDEX idx_venue_scores_venue_id ON venue_performance_scores(venue_id);

-- Index for retrieving scores by week (for platform-wide reporting)
CREATE INDEX idx_venue_scores_week ON venue_performance_scores(score_week_start);

-- Composite index for querying recent scores
CREATE INDEX idx_venue_scores_recent ON venue_performance_scores(venue_id, score_week_start DESC);

COMMENT ON TABLE venue_performance_scores IS 'Weekly venue performance score snapshots with component breakdowns';

```

```
-- =====
-- VENUE PAYOUTS TABLE
-- Tracks revenue and payouts to venues (UI Spec lines 591-593, 736)
-- Supports the platform fee model from PRD Section 2.1
-- =====

CREATE TABLE venue_payouts (
    -- Primary identifier
    id UUID PRIMARY KEY DEFAULT gen_random_uuid(),

    -- Reference to the venue
    venue_id UUID NOT NULL REFERENCES venues(id) ON DELETE CASCADE,

    -- Payout period
    payout_period_start DATE NOT NULL,
    payout_period_end DATE NOT NULL,

    -- Booking metrics for this period
    total_bookings INTEGER DEFAULT 0,
    total_attendees INTEGER DEFAULT 0,

    -- Revenue calculations
    gross_revenue DECIMAL(12, 2) NOT NULL DEFAULT 0 CHECK (gross_revenue >= 0),
    platform_fee_percentage DECIMAL(5, 2) NOT NULL,
    platform_fee_amount DECIMAL(12, 2) NOT NULL DEFAULT 0 CHECK (platform_fee_amount >= 0),
    net_payout DECIMAL(12, 2) NOT NULL DEFAULT 0 CHECK (net_payout >= 0),

    -- Adjustments (refunds, disputes, etc.)
    adjustments_amount DECIMAL(12, 2) DEFAULT 0,
    adjustments_notes TEXT,

    -- Payout status
    status VARCHAR(20) DEFAULT 'pending'
        CHECK (status IN ('pending', 'processing', 'paid', 'failed', 'on_hold')),

    -- Stripe payout tracking
    stripe_payout_id VARCHAR(255),
    stripe_transfer_id VARCHAR(255),

    -- Payment details
    paid_at TIMESTAMP WITH TIME ZONE,
    payment_method VARCHAR(50), -- 'bank_transfer', 'stripe_connect'
    payment_reference VARCHAR(255),

    -- Failure tracking
    failure_reason TEXT,
    retry_count INTEGER DEFAULT 0,
```

```

-- Invoice/statement generation
statement_url VARCHAR(500),
statement_generated_at TIMESTAMP WITH TIME ZONE,

-- Audit timestamps
created_at TIMESTAMP WITH TIME ZONE DEFAULT CURRENT_TIMESTAMP,
updated_at TIMESTAMP WITH TIME ZONE DEFAULT CURRENT_TIMESTAMP,

-- Ensure one payout record per venue per period
CONSTRAINT unique_venue_payout_period UNIQUE (venue_id, payout_period_start)
);

-- Index on venue_id for payout history
CREATE INDEX idx_venue_payouts_venue ON venue_payouts(venue_id);

-- Index on status for processing queue
CREATE INDEX idx_venue_payouts_status ON venue_payouts(status) WHERE status IN ('pending', 'processing');

-- Index on period for reporting
CREATE INDEX idx_venue_payouts_period ON venue_payouts(payout_period_start, payout_period_end);

COMMENT ON TABLE venue_payouts IS 'Venue revenue tracking and payout management';

```

```

-- =====
-- VENUE INTEGRATIONS TABLE
-- Tracks MCP (Model Context Protocol) integrations per venue (PRD Feature 5)
-- Different venues may use different booking systems
-- =====

CREATE TABLE venue_integrations (
    -- Primary identifier
    id UUID PRIMARY KEY DEFAULT gen_random_uuid(),

    -- Reference to the venue
    venue_id UUID NOT NULL REFERENCES venues(id) ON DELETE CASCADE,

    -- Integration type/provider
    integration_type VARCHAR(50) NOT NULL, -- 'mindbody', 'acuity', 'calendly', 'custom_api', 'manual'

    -- Integration display name
    display_name VARCHAR(100),

    -- API configuration (encrypted at app level)
    -- Structure varies by integration type
    api_config JSONB DEFAULT '{}':jsonb,

    -- Authentication credentials reference (stored in secrets manager)
    credentials_secret_id VARCHAR(255),

```

```

-- Webhook configuration for real-time updates
webhook_url VARCHAR(500),
webhook_secret VARCHAR(255),

-- Sync configuration
sync_enabled BOOLEAN DEFAULT true,
sync_frequency_minutes INTEGER DEFAULT 60, -- How often to sync inventory
last_sync_at TIMESTAMP WITH TIME ZONE,
last_sync_status VARCHAR(20), -- 'success', 'partial', 'failed'
last_sync_error TEXT,

-- Inventory sync tracking (PRD Feature 5 - Inventory Sync Watchdog)
last_inventory_check_at TIMESTAMP WITH TIME ZONE,
inventory_discrepancy_count INTEGER DEFAULT 0,

-- Integration health
is_active BOOLEAN DEFAULT true,
health_status VARCHAR(20) DEFAULT 'unknown', -- 'healthy', 'degraded', 'down', 'unknown'
health_check_at TIMESTAMP WITH TIME ZONE,

-- Fallback configuration
fallback_to_email BOOLEAN DEFAULT true, -- Use email if integration fails
fallback_email VARCHAR(255),

-- Audit timestamps
created_at TIMESTAMP WITH TIME ZONE DEFAULT CURRENT_TIMESTAMP,
updated_at TIMESTAMP WITH TIME ZONE DEFAULT CURRENT_TIMESTAMP,

-- One integration type per venue
CONSTRAINT unique_venue_integration UNIQUE (venue_id, integration_type)
);

-- Index on venue_id
CREATE INDEX idx_venue_integrations_venue ON venue_integrations(venue_id);
-- Index on active integrations needing sync
CREATE INDEX idx_venue_integrations_sync ON venue_integrations(last_sync_at, sync_frequency_minutes)
    WHERE is_active = true AND sync_enabled = true;
-- Index on health status for monitoring
CREATE INDEX idx_venue_integrations_health ON venue_integrations(health_status)
    WHERE is_active = true;

COMMENT ON TABLE venue_integrations IS 'MCP integration configurations per venue for booking system sync';

-- =====
-- ACTIVITIES TABLE

```

```
-- Individual activity types offered by venues (e.g., "Robotics Workshop",
-- "Watercolor Painting Class"). Each activity can have multiple scheduled sessions.
```

```
-- =====
```

```
CREATE TABLE activities (
```

```
-- Primary identifier
```

```
id UUID PRIMARY KEY DEFAULT gen_random_uuid(),
```

```
-- Reference to the venue offering this activity
```

```
venue_id UUID NOT NULL REFERENCES venues(id) ON DELETE CASCADE,
```

```
-- Activity name as displayed to parents
```

```
name VARCHAR(255) NOT NULL,
```

```
-- URL-friendly slug for activity detail pages
```

```
slug VARCHAR(255) NOT NULL,
```

```
-- Primary category for filtering and recommendations
```

```
category activity_category NOT NULL,
```

```
-- Brief description for activity cards (150 char limit)
```

```
short_description VARCHAR(150),
```

```
-- Full activity description with learning outcomes
```

```
full_description TEXT,
```

```
-- Learning outcomes as JSONB array for "What your child will learn" section
```

```
-- Structure: ["Learn basic coding concepts", "Build a working robot", ...]
```

```
learning_outcomes JSONB DEFAULT '[]'::jsonb,
```

```
-- Items participants should bring
```

```
-- Structure: ["Comfortable clothes", "Water bottle", ...]
```

```
what_to_bring JSONB DEFAULT '[]'::jsonb,
```

```
-- Age range constraints for eligibility
```

```
min_age INTEGER NOT NULL CHECK (min_age >= 0 AND min_age <= 18),
```

```
max_age INTEGER NOT NULL CHECK (max_age >= 0 AND max_age <= 18),
```

```
-- Duration of each session in minutes
```

```
duration_minutes INTEGER NOT NULL CHECK (duration_minutes > 0),
```

```
-- Credits required to book this activity (deducted at booking time)
```

```
credits_required INTEGER NOT NULL CHECK (credits_required > 0),
```

```
-- Maximum participants per session
```

```
capacity_per_session INTEGER NOT NULL CHECK (capacity_per_session > 0),
```

```
-- Media assets for activity display
```

```
primary_image_url VARCHAR(500),
gallery_images JSONB DEFAULT '[]'::jsonb,

-- Activity characteristics for filtering (stored as JSONB for flexibility)
-- Structure: { "indoor": true, "messy": false, "competitive": false, ... }
activity_tags JSONB DEFAULT '{}'::jsonb,

-- Accessibility features specific to this activity
accessibility_features JSONB DEFAULT '[]'::jsonb,

-- Whether parent must attend with child
parent_attendance_required BOOLEAN DEFAULT false,

-- Skill level requirement
skill_level VARCHAR(20) DEFAULT 'beginner' CHECK (skill_level IN ('beginner', 'intermediate', 'advanced')),

-- Prerequisites text (free-form)
prerequisites TEXT,

-- Activity visibility and status
is_active BOOLEAN DEFAULT true,
is_published BOOLEAN DEFAULT false,

-- =====
-- PINECONE INTEGRATION FIELDS
-- Vector storage is in Pinecone; PostgreSQL only tracks sync status
-- =====

-- Pinecone vector ID for this activity's description embedding
-- Format: "act_{uuid}" - stored in 'activity-embeddings' namespace
pinecone_vector_id VARCHAR(255),

-- Timestamp when embedding was last synced to Pinecone
-- Used by sync job to detect stale vectors needing re-embedding
pinecone_synced_at TIMESTAMP WITH TIME ZONE,

-- Flag indicating if activity content changed since last Pinecone sync
-- Set to TRUE on any update to name, description, tags, or category
-- Sync job resets to FALSE after successful embedding update
pinecone_sync_required BOOLEAN DEFAULT true,

-- Embedding model version used (for bulk re-embedding on model upgrades)
-- Example: "text-embedding-3-small-v1"
embedding_model_version VARCHAR(50),

-- =====
-- DENORMALIZED FIELDS (Computed, not for Pinecone)
```

```

-- =====

-- Average rating computed from reviews (denormalized for performance)
average_rating DECIMAL(3, 2) DEFAULT 0.00 CHECK (average_rating BETWEEN 0 AND 5),
total_reviews INTEGER DEFAULT 0,

-- Total bookings count (denormalized for sorting)
total_bookings INTEGER DEFAULT 0,

-- Audit timestamps
created_at TIMESTAMP WITH TIME ZONE DEFAULT CURRENT_TIMESTAMP,
updated_at TIMESTAMP WITH TIME ZONE DEFAULT CURRENT_TIMESTAMP,
published_at TIMESTAMP WITH TIME ZONE
);

-- Ensure unique slug per venue
ALTER TABLE activities ADD CONSTRAINT unique_activity_slug_per_venue UNIQUE (venue_id, slug);

-- Index on venue_id for retrieving all activities at a venue
CREATE INDEX idx_activities_venue_id ON activities(venue_id);

-- Index on category for filter queries
CREATE INDEX idx_activities_category ON activities(category);

-- Index on age range for eligibility filtering
CREATE INDEX idx_activities_age_range ON activities(min_age, max_age);

-- Index on credits for sorting by price
CREATE INDEX idx_activities_credits ON activities(credits_required);

-- Index on active and published activities
CREATE INDEX idx_activities_active ON activities(is_active, is_published) WHERE is_active = true AND is_published = true;

-- GIN index on activity_tags for JSONB queries (e.g., finding non-messy activities)
CREATE INDEX idx_activities_tags ON activities USING GIN (activity_tags);

-- Index on average rating for sorting
CREATE INDEX idx_activities_rating ON activities(average_rating DESC);

-- Index for Pinecone sync job (find activities needing re-embedding)
CREATE INDEX idx_activities_pinecone_sync ON activities(pinecone_sync_required)
    WHERE pinecone_sync_required = true AND is_active = true;

-- Full-text search index on name and description (FALLBACK when Pinecone unavailable)
-- Per PRD Section 6.3: Fallback to PostgreSQL-based filtering with degraded personalization
CREATE INDEX idx_activities_search ON activities USING GIN (
    to_tsvector('english', name || ' ' || COALESCE(short_description, '') || ' ' || COALESCE(full_description, ''))
);

COMMENT ON TABLE activities IS 'Activity types offered by venues with metadata for search etc';

-- =====

-- ACTIVITY RECURRING SCHEDULES TABLE
-- Defines recurring patterns for activities (UI Spec lines 639-644)

```

```
-- Used to auto-generate activity_sessions based on pattern rules
-- =====

CREATE TABLE activity_recurring_schedules (
    -- Primary identifier
    id UUID PRIMARY KEY DEFAULT gen_random_uuid(),
    -- Reference to the activity this schedule is for
    activity_id UUID NOT NULL REFERENCES activities(id) ON DELETE CASCADE,
    -- Schedule name for venue admin reference
    schedule_name VARCHAR(100),
    -- Days of week this activity runs (stored as array)
    -- Structure: [0, 1, 2, 3, 4, 5, 6] where 0=Sunday, 6=Saturday
    days_of_week JSONB NOT NULL DEFAULT '[]'::jsonb,
    -- Start time for sessions on these days
    start_time TIME NOT NULL,
    -- Duration in minutes (to calculate end_time)
    duration_minutes INTEGER NOT NULL CHECK (duration_minutes > 0),
    -- Capacity per session
    capacity INTEGER NOT NULL CHECK (capacity > 0),
    -- Date range for this recurring schedule
    effective_from DATE NOT NULL,
    effective_until DATE, -- NULL means ongoing
    -- Block booking allocation per session
    platform_allocated_spots INTEGER DEFAULT 0,
    -- Instructor assignment (if consistent)
    default_instructor_name VARCHAR(255),
    -- Is this schedule currently active?
    is_active BOOLEAN DEFAULT true,
    -- Last date sessions were generated up to
    sessions_generated_until DATE,
    -- Audit timestamps
    created_at TIMESTAMP WITH TIME ZONE DEFAULT CURRENT_TIMESTAMP,
    updated_at TIMESTAMP WITH TIME ZONE DEFAULT CURRENT_TIMESTAMP
);
-- Index on activity_id for retrieving schedules
```

```

CREATE INDEX idx_recurring_schedules_activity ON activity_recurring_schedules(activity_id);
-- Index on active schedules for session generation job

CREATE INDEX idx_recurring_schedules_active ON activity_recurring_schedules(is_active, effective_from)
    WHERE is_active = true;
-- GIN index on days_of_week for day filtering

CREATE INDEX idx_recurring_schedules_days ON activity_recurring_schedules USING GIN (days_of_week);

COMMENT ON TABLE activity_recurring_schedules IS 'Recurring schedule patterns for auto-generating activity sessions';

-- =====
-- ACTIVITY SESSIONS TABLE
-- Specific scheduled instances of activities (time slots)
-- Supports the block booking model where platform pre-books slots 30 days ahead
-- =====

CREATE TABLE activity_sessions (
    -- Primary identifier
    id UUID PRIMARY KEY DEFAULT gen_random_uuid(),

    -- Reference to the activity type
    activity_id UUID NOT NULL REFERENCES activities(id) ON DELETE CASCADE,

    -- Reference to recurring schedule (if generated from a pattern)
    recurring_schedule_id UUID REFERENCES activity_recurring_schedules(id) ON DELETE SET NULL,

    -- Session date and time
    session_date DATE NOT NULL,
    start_time TIME NOT NULL,
    end_time TIME NOT NULL,

    -- Timezone for this session (inherited from venue but can be overridden)
    timezone VARCHAR(50) DEFAULT 'America/New_York',

    -- Time of day classification for filtering (UI Spec line 196)
    -- Computed based on start_time: morning (<12), afternoon (12-17), evening (>17)
    time_of_day time_of_day,

    -- Capacity management
    total_capacity INTEGER NOT NULL CHECK (total_capacity > 0),
    booked_count INTEGER DEFAULT 0 CHECK (booked_count >= 0),

    -- Block booking allocation (spots pre-booked by platform)
    platform_allocated_spots INTEGER DEFAULT 0,

    -- Remaining spots for waitlist threshold
    -- Computed as: total_capacity - booked_count
    -- When this reaches 0, new bookings go to waitlist

```

```

-- Session status
is_cancelled BOOLEAN DEFAULT false,
cancellation_reason TEXT,
cancelled_at TIMESTAMP WITH TIME ZONE,
cancelled_by UUID REFERENCES users(id),

-- Session completion tracking
is_completed BOOLEAN DEFAULT false,
completed_at TIMESTAMP WITH TIME ZONE,

-- Instructor assignment (optional tracking)
instructor_name VARCHAR(255),
instructor_notes TEXT,

-- Special notes for this session
session_notes TEXT,

-- Release tracking (72-hour release clause)
inventory_released BOOLEAN DEFAULT false,
inventory_released_at TIMESTAMP WITH TIME ZONE,

-- Audit timestamps
created_at TIMESTAMP WITH TIME ZONE DEFAULT CURRENT_TIMESTAMP,
updated_at TIMESTAMP WITH TIME ZONE DEFAULT CURRENT_TIMESTAMP,

-- Ensure end time is after start time
CONSTRAINT valid_session_times CHECK (end_time > start_time),
-- Ensure booked count doesn't exceed capacity
CONSTRAINT valid_booking_count CHECK (booked_count <= total_capacity)
);

-- Index on activity_id for retrieving all sessions of an activity
CREATE INDEX idx_sessions_activity_id ON activity_sessions(activity_id);
-- Index on session_date for date range queries
CREATE INDEX idx_sessions_date ON activity_sessions(session_date);
-- Composite index for finding available sessions (most common query)
CREATE INDEX idx_sessions_availability ON activity_sessions(activity_id, session_date, is_cancelled)
    WHERE is_cancelled = false;
-- Index for finding sessions with available spots
CREATE INDEX idx_sessions_available_spots ON activity_sessions(session_date, booked_count, total_capacity)
    WHERE is_cancelled = false AND booked_count < total_capacity;
-- Index on cancelled status for reporting
CREATE INDEX idx_sessions_cancelled ON activity_sessions(is_cancelled) WHERE is_cancelled = true;
-- Index on time_of_day for filtering (UI Spec line 196)
CREATE INDEX idx_sessions_time_of_day ON activity_sessions(time_of_day, session_date)
    WHERE is_cancelled = false;

```

COMMENT ON TABLE activity_sessions **IS** 'Scheduled time slots for activities with capacity and booking management';

2.5 Children and Profile Tables

sql

```

-- =====
-- CHILDREN TABLE
-- Child profiles registered by parents for activity matching
-- Subject to COPPA/GDPR-K compliance with strict data isolation
-- =====

CREATE TABLE children (
    -- Primary identifier (never exposed externally, use initials for privacy)
    id UUID PRIMARY KEY DEFAULT gen_random_uuid(),

    -- Reference to the parent user who registered this child
    parent_id UUID NOT NULL REFERENCES users(id) ON DELETE CASCADE,

    -- Child's first name (no last name stored for privacy)
    first_name VARCHAR(100) NOT NULL,

    -- Date of birth for age calculation and eligibility checks
    date_of_birth DATE NOT NULL,

    -- Optional gender for personalization (not required)
    gender VARCHAR(20) CHECK (gender IN ('boy', 'girl', 'non_binary', 'prefer_not_to_say')),

    -- Profile photo URL (parent-uploaded)
    avatar_url VARCHAR(500),

    -- Onboarding quiz responses (3-question quiz per PRD)
    energy_level energy_level,
    social_preference social_preference,

    -- Interest categories selected during onboarding (up to 3)
    -- Stored as JSONB array: ["stem", "arts_crafts", "building"]
    interests JSONB DEFAULT '[]'::jsonb,

    -- Activities to avoid (parent-specified constraints)
    -- Structure: ["competitive", "messy", "loud", "physical_contact"]
    activities_to_avoid JSONB DEFAULT '[]'::jsonb,

    -- Medical and allergy notes (shared with venues for safety)
    allergies_medical_notes TEXT,

    -- Accessibility needs
    -- Structure: ["wheelchair_accessible", "sensory_friendly", ...]
    accessibility_needs JSONB DEFAULT '[]'::jsonb,

-- =====
-- PINECONE INTEGRATION FIELDS
-- Child preference vectors stored in Pinecone 'child-preferences' namespace

```

-- Used for semantic matching with activities and sibling compatibility

-- =====

-- Pinecone vector ID for this child's preference embedding

-- Format: "child_{uuid}" - combines quiz answers + feedback history

pinecone_vector_id VARCHAR(255),

-- Timestamp when preference vector was last synced to Pinecone

pinecone_synced_at TIMESTAMP WITH TIME ZONE,

-- Flag indicating preferences changed since last Pinecone sync

-- Set TRUE when: quiz answers change, new feedback received, interests updated

pinecone_sync_required BOOLEAN DEFAULT true,

-- Embedding model version used for this child's vector

embedding_model_version VARCHAR(50),

-- Count of feedback events incorporated into current vector

-- Helps track vector evolution over time

feedback_incorporated_count INTEGER DEFAULT 0,

-- =====

-- ACTIVITY TRACKING

-- =====

-- Denormalized activity count for this month (for dashboard display)

activities_this_month INTEGER DEFAULT 0,

-- =====

-- COPPA COMPLIANCE FIELDS

-- =====

-- Parental consent tracking (COPPA compliance)

parental_consent_given BOOLEAN DEFAULT false,

parental_consent_date TIMESTAMP WITH TIME ZONE,

parental_consent_ip INET,

-- Soft delete support for data retention compliance

is_active BOOLEAN DEFAULT true,

deactivated_at TIMESTAMP WITH TIME ZONE,

-- Data deletion request tracking (GDPR-K right to deletion)

-- NOTE: When deletion occurs, must also delete from Pinecone

deletion_requested BOOLEAN DEFAULT false,

deletion_requested_at TIMESTAMP WITH TIME ZONE,

deletion_scheduled_for DATE,

```

-- Audit timestamps
created_at TIMESTAMP WITH TIME ZONE DEFAULT CURRENT_TIMESTAMP,
updated_at TIMESTAMP WITH TIME ZONE DEFAULT CURRENT_TIMESTAMP
);

-- Index on parent_id for retrieving all children of a parent
CREATE INDEX idx_children_parent_id ON children(parent_id);
-- Index on date_of_birth for age calculation queries
CREATE INDEX idx_children_dob ON children(date_of_birth);
-- Index on active children
CREATE INDEX idx_children_active ON children(parent_id, is_active) WHERE is_active = true;
-- GIN index on interests for filtering children by interest
CREATE INDEX idx_children_interests ON children USING GIN (interests);
-- Index for deletion queue processing
CREATE INDEX idx_children_deletion_pending ON children(deletion_scheduled_for)
    WHERE deletion_requested = true;
-- Index for Pinecone sync job (find children needing vector update)
CREATE INDEX idx_children_pinecone_sync ON children(pinecone_sync_required)
    WHERE pinecone_sync_required = true AND is_active = true;

COMMENT ON TABLE children IS 'Child profiles for activity matching, subject to COPPA/GDPR-K compliance. Preferenc';

```

```

-- =====
-- CHILD FAVORITES TABLE
-- Activities favorited/saved by parents for quick access
-- =====

CREATE TABLE childFavorites (
    -- Primary identifier
    id UUID PRIMARY KEY DEFAULT gen_random_uuid(),
    -- Reference to the child this favorite is for
    child_id UUID NOT NULL REFERENCES children(id) ON DELETE CASCADE,
    -- Reference to the favorited activity
    activity_id UUID NOT NULL REFERENCES activities(id) ON DELETE CASCADE,
    -- Timestamp when favorited
    created_at TIMESTAMP WITH TIME ZONE DEFAULT CURRENT_TIMESTAMP,
    -- Ensure a child can only favorite an activity once
    CONSTRAINT unique_child_favorite UNIQUE (child_id, activity_id)
);

-- Index for retrieving favorites for a child
CREATE INDEX idx_favorites_child_id ON childFavorites(child_id);

```

COMMENT ON TABLE childFavorites IS 'Parent-saved favorite activities for quick access';

2.6 Subscription and Credit Management Tables

sql

```
-- =====
-- CORPORATE SUBSCRIPTIONS TABLE
-- Subscription periods for corporate accounts with credit allocation rules
-- =====

CREATE TABLE corporate_subscriptions (
    -- Primary identifier
    id UUID PRIMARY KEY DEFAULT gen_random_uuid(),

    -- Reference to the corporate account
    corporate_id UUID NOT NULL REFERENCES corporates(id) ON DELETE CASCADE,

    -- Subscription tier determining features and credit limits
    tier subscription_tier NOT NULL,

    -- Number of employees covered by this subscription
    covered_employees INTEGER NOT NULL CHECK (covered_employees > 0),

    -- Credits allocated per employee per month
    credits_per_employee INTEGER NOT NULL CHECK (credits_per_employee > 0),

    -- Monthly subscription cost (for invoicing)
    monthly_cost DECIMAL(12, 2) NOT NULL CHECK (monthly_cost >= 0),

    -- Billing cycle day of month (1-28 to avoid month-end issues)
    billing_cycle_day INTEGER DEFAULT 1 CHECK (billing_cycle_day BETWEEN 1 AND 28),

    -- Subscription period
    start_date DATE NOT NULL,
    end_date DATE,

    -- Auto-renewal configuration
    auto_renew BOOLEAN DEFAULT true,

    -- Premium venue access flag (for premium/enterprise tiers)
    premium_venue_access BOOLEAN DEFAULT false,

    -- Custom terms for enterprise accounts (stored as JSONB)
    custom_terms JSONB,

    -- Subscription status
    is_active BOOLEAN DEFAULT true,
    cancelled_at TIMESTAMP WITH TIME ZONE,
    cancellation_reason TEXT,

    -- Audit timestamps
    created_at TIMESTAMP WITH TIME ZONE DEFAULT CURRENT_TIMESTAMP,
```

```

updated_at TIMESTAMP WITH TIME ZONE DEFAULT CURRENT_TIMESTAMP
);

-- Index on corporate_id for retrieving subscription history
CREATE INDEX idx_subscriptions_corporate_id ON corporate_subscriptions(corporate_id);
-- Index on active subscriptions
CREATE INDEX idx_subscriptions_active ON corporate_subscriptions(corporate_id, is_active)
    WHERE is_active = true;
-- Index on end dates for renewal processing
CREATE INDEX idx_subscriptions_end_date ON corporate_subscriptions(end_date);

COMMENT ON TABLE corporate_subscriptions IS 'Corporate subscription periods with credit allocation rules';

-- =====
-- EMPLOYEE INVITATIONS TABLE
-- Tracks HR-initiated invitations to employees (UI Spec lines 506-513)
-- Supports bulk CSV upload and manual invitation workflows
-- =====

CREATE TABLE employee_invitations (
    -- Primary identifier
    id UUID PRIMARY KEY DEFAULT gen_random_uuid(),

    -- Reference to the corporate account
    corporate_id UUID NOT NULL REFERENCES corporates(id) ON DELETE CASCADE,

    -- Reference to the HR admin who sent the invitation
    invited_by UUID NOT NULL REFERENCES users(id) ON DELETE CASCADE,

    -- Invitee details
    email VARCHAR(255) NOT NULL,
    first_name VARCHAR(100),
    last_name VARCHAR(100),
    department VARCHAR(100),

    -- Credits to allocate once they register
    credits_to_allocate INTEGER NOT NULL CHECK (credits_to_allocate > 0),

    -- Invitation token for registration link
    invitation_token VARCHAR(255) NOT NULL UNIQUE,

    -- Invitation status tracking
    status invitation_status DEFAULT 'pending',

    -- Token expiration (typically 7-14 days)
    expires_at TIMESTAMP WITH TIME ZONE NOT NULL,
);

```

```

-- If accepted, reference to the created user account
accepted_user_id UUID REFERENCES users(id) ON DELETE SET NULL,
accepted_at TIMESTAMP WITH TIME ZONE,

-- Resend tracking
resend_count INTEGER DEFAULT 0,
last_resent_at TIMESTAMP WITH TIME ZONE,

-- Revocation details
revoked_at TIMESTAMP WITH TIME ZONE,
revoked_by UUID REFERENCES users(id),
revoke_reason TEXT,

-- Batch import tracking (for CSV uploads)
batch_import_id UUID, -- Groups invitations from same CSV upload

-- Audit timestamps
created_at TIMESTAMP WITH TIME ZONE DEFAULT CURRENT_TIMESTAMP,
updated_at TIMESTAMP WITH TIME ZONE DEFAULT CURRENT_TIMESTAMP
);

-- Index on corporate_id for listing all invitations
CREATE INDEX idx_invitations_corporate_id ON employee_invitations(corporate_id);
-- Index on email for duplicate checking
CREATE INDEX idx_invitations_email ON employee_invitations(corporate_id, email);
-- Index on status for filtering pending invitations
CREATE INDEX idx_invitations_status ON employee_invitations(status);
-- Index on token for registration lookup
CREATE INDEX idx_invitations_token ON employee_invitations(invitation_token) WHERE status = 'pending';
-- Index on expiration for cleanup job
CREATE INDEX idx_invitations_expiry ON employee_invitations(expires_at) WHERE status = 'pending';
-- Index on batch for grouping CSV imports
CREATE INDEX idx_invitations_batch ON employee_invitations(batch_import_id) WHERE batch_import_id IS NOT NULL;

COMMENT ON TABLE employee_invitations IS 'HR-initiated employee invitations with token-based registration';

```

```

-- =====
-- USER CREDIT BALANCES TABLE
-- Current credit balance for each parent user (denormalized from ledger)
-- Reset monthly based on corporate subscription billing cycle
-- =====

CREATE TABLE user_credit_balances (
    -- Primary identifier
    id UUID PRIMARY KEY DEFAULT gen_random_uuid(),
    -- Reference to the parent user

```

```

user_id UUID NOT NULL REFERENCES users(id) ON DELETE CASCADE,
-- Current billing period (month start date)
billing_period_start DATE NOT NULL,
billing_period_end DATE NOT NULL,
-- Credit allocation for this period
credits_allocated INTEGER NOT NULL CHECK (credits_allocated >= 0),
-- Current available balance (allocated - used + refunded)
credits_available INTEGER NOT NULL CHECK (credits_available >= 0),
-- Credits already used this period
credits_used INTEGER DEFAULT 0 CHECK (credits_used >= 0),
-- Credits refunded this period
credits_refunded INTEGER DEFAULT 0 CHECK (credits_refunded >= 0),
-- Credits forfeited (no-shows and late cancellations)
credits_forfeited INTEGER DEFAULT 0 CHECK (credits_forfeited >= 0),
-- Timestamp of last credit refresh
last_refresh_at TIMESTAMP WITH TIME ZONE,
-- Expiry warning sent flag (7 days before period end)
expiry_warning_sent BOOLEAN DEFAULT false,
expiry_warning_sent_at TIMESTAMP WITH TIME ZONE,
-- Audit timestamps
created_at TIMESTAMP WITH TIME ZONE DEFAULT CURRENT_TIMESTAMP,
updated_at TIMESTAMP WITH TIME ZONE DEFAULT CURRENT_TIMESTAMP,
-- Ensure one active balance record per user per billing period
CONSTRAINT unique_user_billing_period UNIQUE (user_id, billing_period_start)
);
-- Index on user_id for retrieving current balance
CREATE INDEX idx_credit_balances_user_id ON user_credit_balances(user_id);
-- Index for finding expiring balances (7-day warning)
CREATE INDEX idx_credit_balances_expiry ON user_credit_balances(billing_period_end, expiry_warning_sent)
WHERE expiry_warning_sent = false;
COMMENT ON TABLE user_credit_balances IS 'Current credit balance per parent per billing period';
-- =====
-- CREDIT LEDGER TABLE

```

```
-- Immutable transaction log of all credit movements for audit trail
-- Supports the financial operations described in PRD Section 2.2
-- =====

CREATE TABLE credit_ledger (
    -- Primary identifier
    id UUID PRIMARY KEY DEFAULT gen_random_uuid(),

    -- Reference to the user whose credits are affected
    user_id UUID NOT NULL REFERENCES users(id) ON DELETE CASCADE,

    -- Reference to the balance record being modified
    balance_id UUID NOT NULL REFERENCES user_credit_balances(id) ON DELETE CASCADE,

    -- Type of credit transaction
    transaction_type credit_transaction_type NOT NULL,

    -- Credit amount (positive for credits in, negative for credits out)
    amount INTEGER NOT NULL,

    -- Running balance after this transaction
    balance_after INTEGER NOT NULL,

    -- Optional reference to related booking (for booking debits/refunds)
    booking_id UUID, -- FK added after bookings table creation

    -- Description of the transaction
    description TEXT,

    -- Reference to admin who made adjustment (for manual adjustments)
    adjusted_by UUID REFERENCES users(id),

    -- Idempotency key to prevent duplicate transactions
    idempotency_key VARCHAR(255),

    -- Transaction timestamp (immutable)
    created_at TIMESTAMP WITH TIME ZONE DEFAULT CURRENT_TIMESTAMP,

    -- Ensure idempotency key is unique if provided
    CONSTRAINT unique_idempotency_key UNIQUE (idempotency_key)
);

-- Index on user_id for retrieving transaction history
CREATE INDEX idx_credit_ledger_user_id ON credit_ledger(user_id);
-- Index on balance_id for period-specific queries
CREATE INDEX idx_credit_ledger_balance_id ON credit_ledger(balance_id);
-- Index on booking_id for finding credit transactions for a booking
CREATE INDEX idx_credit_ledger_booking_id ON credit_ledger(booking_id) WHERE booking_id IS NOT NULL;
```

-- Index on transaction type for reporting

CREATE INDEX idx_credit_ledger_type **ON** credit_ledger(transaction_type);

-- Index on created_at for date range queries

CREATE INDEX idx_credit_ledger_created_at **ON** credit_ledger(created_at);

COMMENT ON TABLE credit_ledger **IS** 'Immutable audit log of all credit transactions';

2.7 Booking and Waitlist Tables

sql

```

-- =====
-- BOOKINGS TABLE
-- Activity reservations made by parents for their children
-- Core transaction table with status tracking through complete lifecycle
-- =====

CREATE TABLE bookings (
    -- Primary identifier
    id UUID PRIMARY KEY DEFAULT gen_random_uuid(),

    -- Reference to the booked activity session
    session_id UUID NOT NULL REFERENCES activity_sessions(id) ON DELETE RESTRICT,

    -- Reference to the child attending (for single-child bookings)
    child_id UUID NOT NULL REFERENCES children(id) ON DELETE RESTRICT,

    -- Reference to the parent who made the booking
    booked_by UUID NOT NULL REFERENCES users(id) ON DELETE RESTRICT,

    -- Booking status tracking lifecycle
    status booking_status DEFAULT 'pending',

    -- Credits charged for this booking
    credits_charged INTEGER NOT NULL CHECK (credits_charged > 0),

    -- Group booking reference (links sibling bookings together)
    -- All bookings made in a single group transaction share this ID
    group_booking_id UUID,

    -- Confirmation details
    confirmation_code VARCHAR(20) NOT NULL UNIQUE,
    confirmed_at TIMESTAMP WITH TIME ZONE,

    -- Cancellation details (if applicable)
    cancelled_at TIMESTAMP WITH TIME ZONE,
    cancellation_reason cancellation_reason_type, -- Structured reason (UI Spec line 330)
    cancellation_notes TEXT, -- Additional details if reason is 'other'

    -- Refund tracking
    refund_eligible BOOLEAN,
    refund_processed BOOLEAN DEFAULT false,
    refund_processed_at TIMESTAMP WITH TIME ZONE,
    credits_refunded INTEGER DEFAULT 0,

    -- Attendance tracking (marked by venue)
    attendance_marked BOOLEAN DEFAULT false,
    attendance_marked_at TIMESTAMP WITH TIME ZONE,

```

```
attendance_marked_by UUID REFERENCES users(id),  
  
-- Special notes from parent (allergies reminder, etc.)  
parent_notes TEXT,  
  
-- Venue-side notes for this booking  
venue_notes TEXT,  
  
-- AI curation tracking (was this from a suggestion?)  
from_ai_suggestion BOOLEAN DEFAULT false,  
ai_suggestion_id UUID, -- FK added after ai_suggestions table  
  
-- Source tracking for analytics  
booking_source VARCHAR(50) DEFAULT 'browse', -- 'browse', 'suggestion', 'search', 'waitlist', 'rebook'  
  
-- Reference to original booking if this is a "Book Again" (UI Spec line 311)  
original_booking_id UUID REFERENCES bookings(id) ON DELETE SET NULL,  
  
-- Calendar event tracking  
calendar_event_sent BOOLEAN DEFAULT false,  
  
-- Reminder notifications sent  
reminder_24h_sent BOOLEAN DEFAULT false,  
reminder_1h_sent BOOLEAN DEFAULT false,  
  
-- Audit timestamps  
created_at TIMESTAMP WITH TIME ZONE DEFAULT CURRENT_TIMESTAMP,  
updated_at TIMESTAMP WITH TIME ZONE DEFAULT CURRENT_TIMESTAMP,  
  
-- Row-level locking version for optimistic concurrency control (per PRD Section 8)  
version INTEGER DEFAULT 1  
);  
  
-- Index on session_id for finding all bookings for a session (attendance list)  
CREATE INDEX idx_bookings_session_id ON bookings(session_id);  
-- Index on child_id for finding all bookings for a child  
CREATE INDEX idx_bookings_child_id ON bookings(child_id);  
-- Index on booked_by for parent's booking history  
CREATE INDEX idx_bookings_booked_by ON bookings(booked_by);  
-- Index on status for filtering by booking state  
CREATE INDEX idx_bookings_status ON bookings(status);  
-- Index on confirmation_code for lookup  
CREATE INDEX idx_bookings_confirmation ON bookings(confirmation_code);  
-- Index on group_booking_id for finding sibling bookings  
CREATE INDEX idx_bookings_group ON bookings(group_booking_id) WHERE group_booking_id IS NOT NULL;  
-- Composite index for upcoming bookings (dashboard query)  
CREATE INDEX idx_bookings_upcoming ON bookings(booked_by, status, created_at DESC)
```

```

WHERE status IN ('pending', 'confirmed');

-- Index for AI suggestion tracking
CREATE INDEX idx_bookings_ai_suggestion ON bookings(ai_suggestion_id) WHERE ai_suggestion_id IS NOT NULL;

-- Add foreign key from credit_ledger to bookings
ALTER TABLE credit_ledger ADD CONSTRAINT fk_credit_ledger_booking
FOREIGN KEY (booking_id) REFERENCES bookings(id) ON DELETE SET NULL;

COMMENT ON TABLE bookings IS 'Activity reservations with full lifecycle tracking and group booking support';

-- =====
-- WAITLIST TABLE
-- Queue for full activity sessions with automatic notification
-- Priority is first-come-first-served with 4-hour confirmation window
-- =====

CREATE TABLE waitlist (
    -- Primary identifier
    id UUID PRIMARY KEY DEFAULT gen_random_uuid(),

    -- Reference to the activity session they're waiting for
    session_id UUID NOT NULL REFERENCES activity_sessions(id) ON DELETE CASCADE,

    -- Reference to the child wanting to attend
    child_id UUID NOT NULL REFERENCES children(id) ON DELETE CASCADE,

    -- Reference to the parent on the waitlist
    user_id UUID NOT NULL REFERENCES users(id) ON DELETE CASCADE,

    -- Position in the waitlist queue (auto-incremented per session)
    position INTEGER NOT NULL,

    -- Timestamp when joined waitlist (for FIFO ordering)
    joined_at TIMESTAMP WITH TIME ZONE DEFAULT CURRENT_TIMESTAMP,

    -- Notification tracking when spot opens
    spot_offered_at TIMESTAMP WITH TIME ZONE,
    spot_expires_at TIMESTAMP WITH TIME ZONE, -- 4 hours after offered

    -- Response tracking
    response VARCHAR(20) CHECK (response IN ('accepted', 'declined', 'expired')),
    responded_at TIMESTAMP WITH TIME ZONE,

    -- If accepted, reference to the created booking
    converted_booking_id UUID REFERENCES bookings(id),

    -- Active status (removed from waitlist or not)
);

```

```

is_active BOOLEAN DEFAULT true,
removed_at TIMESTAMP WITH TIME ZONE,
removed_reason VARCHAR(50), -- 'user_removed', 'spot_accepted', 'spot_declined', 'spot_expired'

-- Audit timestamps
created_at TIMESTAMP WITH TIME ZONE DEFAULT CURRENT_TIMESTAMP,
updated_at TIMESTAMP WITH TIME ZONE DEFAULT CURRENT_TIMESTAMP,

-- Ensure a child can only be on waitlist once per session
CONSTRAINT unique_waitlist_entry UNIQUE (session_id, child_id)
);

-- Index on session_id for finding waitlist for a session
CREATE INDEX idx_waitlist_session_id ON waitlist(session_id);
-- Index on user_id for parent's waitlist view
CREATE INDEX idx_waitlist_user_id ON waitlist(user_id);
-- Index on child_id for finding child's waitlist entries
CREATE INDEX idx_waitlist_child_id ON waitlist(child_id);
-- Index for finding active waitlist entries ordered by position
CREATE INDEX idx_waitlist_active ON waitlist(session_id, position) WHERE is_active = true;
-- Index for finding entries with pending spot offers (for expiration processing)
CREATE INDEX idx_waitlist_pending_offers ON waitlist(spot_expires_at)
    WHERE spot_offered_at IS NOT NULL AND response IS NULL;

COMMENT ON TABLE waitlist IS 'FIFO queue for full sessions with 4-hour confirmation window';

```

2.8 Feedback, Reviews, and Notifications Tables

sql

```

-- =====
-- BOOKING FEEDBACK TABLE
-- Quick thumbs up/down feedback collected 1 hour post-activity
-- Used for real-time vector updates and venue scoring
-- =====

CREATE TABLE booking_feedback (
    -- Primary identifier
    id UUID PRIMARY KEY DEFAULT gen_random_uuid(),

    -- Reference to the booking this feedback is for (one feedback per booking)
    booking_id UUID NOT NULL UNIQUE REFERENCES bookings(id) ON DELETE CASCADE,

    -- Simple sentiment: true = thumbs up, false = thumbs down
    is_positive BOOLEAN NOT NULL,

    -- Optional tags for positive feedback
    -- Structure: ["fun", "educational", "social", "creative"]
    positive_tags JSONB DEFAULT '[]'::jsonb,

    -- Optional tags for negative feedback
    -- Structure: ["organization", "content", "instructor", "facility"]
    negative_tags JSONB DEFAULT '[]'::jsonb,

    -- Optional text comment
    comment TEXT,

    -- Timestamp when feedback was submitted
    submitted_at TIMESTAMP WITH TIME ZONE DEFAULT CURRENT_TIMESTAMP,

    -- Flag indicating if this feedback has been processed for vector update
    vector_processed BOOLEAN DEFAULT false,
    vector_processed_at TIMESTAMP WITH TIME ZONE
);

-- Index on booking_id (already has unique constraint)
-- Index for unprocessed feedback (for vector update batch job)
CREATE INDEX idx_feedback_unprocessed ON booking_feedback(vector_processed)
    WHERE vector_processed = false;

COMMENT ON TABLE booking_feedback IS 'Quick post-activity sentiment feedback for AI vector updates';

-- =====
-- REVIEWS TABLE
-- Detailed reviews with star ratings for venue performance scoring
-- Aggregated for activity and venue average ratings

```

```
-- =====
CREATE TABLE reviews (
    -- Primary identifier
    id UUID PRIMARY KEY DEFAULT gen_random_uuid(),
    -- Reference to the booking this review is for (one review per booking)
    booking_id UUID NOT NULL UNIQUE REFERENCES bookings(id) ON DELETE CASCADE,
    -- Reference to the activity being reviewed
    activity_id UUID NOT NULL REFERENCES activities(id) ON DELETE CASCADE,
    -- Reference to the venue being reviewed
    venue_id UUID NOT NULL REFERENCES venues(id) ON DELETE CASCADE,
    -- Reference to the user who wrote the review
    user_id UUID NOT NULL REFERENCES users(id) ON DELETE CASCADE,
    -- Star rating (1-5)
    rating INTEGER NOT NULL CHECK (rating BETWEEN 1 AND 5),
    -- Would book again? (for repeat booking metric)
    would_book_again BOOLEAN,
    -- Would recommend venue? (for venue scoring)
    would_recommend_venue BOOLEAN,
    -- Written review text (optional)
    review_text TEXT,
    -- Review visibility (can be hidden by admin if inappropriate)
    is_visible BOOLEAN DEFAULT true,
    hidden_reason TEXT,
    hidden_by UUID REFERENCES users(id),
    hidden_at TIMESTAMP WITH TIME ZONE,
    -- Venue response to review (optional)
    venue_response TEXT,
    venue_responded_at TIMESTAMP WITH TIME ZONE,
    venue_responded_by UUID REFERENCES users(id),
    -- Moderation status
    is_flagged BOOLEAN DEFAULT false,
    flagged_reason TEXT,
    -- Timestamp when review was submitted
    submitted_at TIMESTAMP WITH TIME ZONE DEFAULT CURRENT_TIMESTAMP,
```

```

-- Audit timestamps
created_at TIMESTAMP WITH TIME ZONE DEFAULT CURRENT_TIMESTAMP,
updated_at TIMESTAMP WITH TIME ZONE DEFAULT CURRENT_TIMESTAMP
);

-- Index on activity_id for retrieving reviews for an activity
CREATE INDEX idx_reviews_activity_id ON reviews(activity_id);
-- Index on venue_id for venue performance scoring
CREATE INDEX idx_reviews_venue_id ON reviews(venue_id);
-- Index on user_id for user's review history
CREATE INDEX idx_reviews_user_id ON reviews(user_id);
-- Index on rating for aggregation queries
CREATE INDEX idx_reviews_rating ON reviews(rating);
-- Index on visible reviews for public display
CREATE INDEX idx_reviews_visible ON reviews(activity_id, is_visible, submitted_at DESC)
    WHERE is_visible = true;

```

COMMENT ON TABLE reviews IS 'Detailed reviews with ratings for venue performance scoring';

```

-- =====
-- NOTIFICATIONS TABLE
-- Multi-channel notification delivery tracking
-- Supports email, SMS, WhatsApp, push, and in-app notifications
-- =====

CREATE TABLE notifications (
    -- Primary identifier
    id UUID PRIMARY KEY DEFAULT gen_random_uuid(),
    -- Reference to the recipient user
    user_id UUID NOT NULL REFERENCES users(id) ON DELETE CASCADE,
    -- Notification type for categorization and filtering
    notification_type VARCHAR(50) NOT NULL,
    -- Types: 'booking_confirmed', 'bookingReminder', 'booking_cancelled_venue',
    --       'monthly_suggestions', 'waitlist_spot', 'feedback_request',
    --       'credits_expiring', 'new_activities', 'system_announcement'
    -- Notification priority
    priority notification_priority DEFAULT 'normal',
    -- Notification title/subject
    title VARCHAR(255) NOT NULL,
    -- Notification body content
    body TEXT NOT NULL,

```

```

-- Structured data payload for rich notifications (e.g., booking details)
data_payload JSONB DEFAULT '{}'::jsonb,

-- Delivery channel used
channel notification_channel NOT NULL,

-- Deep link URL for in-app navigation
action_url VARCHAR(500),

-- Related entity references (polymorphic)
related_entity_type VARCHAR(50), -- 'booking', 'activity', 'waitlist', etc.
related_entity_id UUID,

-- Delivery status tracking
is_sent BOOLEAN DEFAULT false,
sent_at TIMESTAMP WITH TIME ZONE,
send_error TEXT,

-- External delivery ID (from email/SMS/push provider)
external_delivery_id VARCHAR(255),

-- Read/seen tracking
is_read BOOLEAN DEFAULT false,
read_at TIMESTAMP WITH TIME ZONE,

-- Scheduled delivery (for future notifications like reminders)
scheduled_for TIMESTAMP WITH TIME ZONE,

-- Audit timestamps
created_at TIMESTAMP WITH TIME ZONE DEFAULT CURRENT_TIMESTAMP
);

-- Index on user_id for retrieving user's notifications
CREATE INDEX idx_notifications_user_id ON notifications(user_id);

-- Index for unread notifications (notification center badge)
CREATE INDEX idx_notifications_unread ON notifications(user_id, is_read, created_at DESC)
    WHERE is_read = false;

-- Index on notification_type for filtering
CREATE INDEX idx_notifications_type ON notifications(notification_type);

-- Index on scheduled notifications for delivery processing
CREATE INDEX idx_notifications_scheduled ON notifications(scheduled_for)
    WHERE is_sent = false AND scheduled_for IS NOT NULL;

-- Index for pending sends
CREATE INDEX idx_notifications_pending ON notifications(is_sent, created_at) WHERE is_sent = false;

```

COMMENT ON TABLE notifications IS 'Multi-channel notification delivery with status tracking';

2.9 AI Curation and Suggestions Tables

sql

```

-- =====
-- AI CURATION SUGGESTIONS TABLE
-- Monthly personalized activity suggestions generated by LangGraph agent
-- Supports individual and sibling group recommendations
-- NOTE: Match scores come from Pinecone vector similarity; this table stores
--       the final suggestions and tracks parent responses for learning
-- =====

CREATE TABLE ai_suggestions (
    -- Primary identifier
    id UUID PRIMARY KEY DEFAULT gen_random_uuid(),
    -- Reference to the parent user receiving suggestions
    user_id UUID NOT NULL REFERENCES users(id) ON DELETE CASCADE,
    -- Billing period these suggestions are for
    suggestion_month DATE NOT NULL, -- First day of the month
    -- Children these suggestions are for (can be single or multiple for siblings)
    -- Stored as JSONB array of child IDs
    child_ids JSONB NOT NULL, -- Structure: ["uuid1", "uuid2"]
    -- Reference to the suggested activity session
    session_id UUID NOT NULL REFERENCES activity_sessions(id) ON DELETE CASCADE,
    -- Reference to the activity (denormalized for easier queries)
    activity_id UUID NOT NULL REFERENCES activities(id) ON DELETE CASCADE,
    -- Suggestion ranking (1, 2, 3 for top 3 suggestions)
    rank INTEGER CHECK (rank BETWEEN 1 AND 10),
    -- Is this a sibling/group suggestion?
    is_group_suggestion BOOLEAN DEFAULT false,
    -- AI explanation of why this was suggested (generated by LLM)
    explanation TEXT,
    -- =====
    -- PINECONE-DERIVED SCORES
    -- These scores come from Pinecone vector similarity search
    -- =====
    -- Overall match score from Pinecone (cosine similarity, 0-1)
    -- For single child: direct child→activity similarity
    -- For siblings: computed intersection score
    match_score DECIMAL(4, 3) CHECK (match_score BETWEEN 0 AND 1),

```

```

-- Individual child fit scores for group suggestions (from Pinecone)
-- Structure: { "child_uuid1": 0.85, "child_uuid2": 0.78 }
individual_fit_scores JSONB,  

-- Pinecone query details for debugging/auditing
pinecone_query_id VARCHAR(255),  

-- ======  

-- PARENT RESPONSE TRACKING
-- ======  

-- Parent response to suggestion
parent_response VARCHAR(20), -- 'accepted', 'rejected', 'ignored'
responded_at TIMESTAMP WITH TIME ZONE,  

-- If accepted, reference to the created booking
converted_booking_id UUID REFERENCES bookings(id),  

-- Rejection reason if declined (used to improve future suggestions)
rejection_reason TEXT,  

-- ======  

-- AUDIT & DEBUGGING
-- ======  

-- LangSmith trace ID for AI decision auditing
langsmith_trace_id VARCHAR(255),  

-- Generation metadata
generated_at TIMESTAMP WITH TIME ZONE DEFAULT CURRENT_TIMESTAMP,  

-- Expiry (suggestions expire at end of month)
expires_at TIMESTAMP WITH TIME ZONE,  

-- Audit timestamps
created_at TIMESTAMP WITH TIME ZONE DEFAULT CURRENT_TIMESTAMP,
updated_at TIMESTAMP WITH TIME ZONE DEFAULT CURRENT_TIMESTAMP
);  

-- Index on user_id for retrieving user's suggestions
CREATE INDEX idxSuggestionsUser_id ON aiSuggestions(user_id);
-- Index on suggestion_month for current month queries
CREATE INDEX idxSuggestionsMonth ON aiSuggestions(suggestion_month);
-- Composite index for parent dashboard (current month suggestions)
CREATE INDEX idxSuggestionsCurrent ON aiSuggestions(user_id, suggestion_month, rank)
    WHERE parent_response IS NULL;
-- Index on activity_id for suggestion analytics

```

```
CREATE INDEX idx_suggestions_activity_id ON ai_suggestions(activity_id);
-- GIN index on child_ids for finding suggestions for specific children
CREATE INDEX idx_suggestions_child_ids ON ai_suggestions USING GIN (child_ids);

COMMENT ON TABLE ai_suggestions IS 'Monthly AI-curated activity suggestions with sibling matching';
```

```
-- =====
-- NATURAL LANGUAGE SEARCH LOGS TABLE
-- Analytics and audit log for NL search queries
-- NOTE: Actual semantic search is performed in Pinecone; this table is for:
-- 1. Tracking search accuracy (90% target per PRD Section 7.3)
-- 2. Analyzing user search patterns for improvements
-- 3. Debugging failed searches
-- 4. Compliance auditing
-- =====

CREATE TABLE nl_search_logs (
    -- Primary identifier
    id UUID PRIMARY KEY DEFAULT gen_random_uuid(),
    -- Reference to the user who made the search
    user_id UUID NOT NULL REFERENCES users(id) ON DELETE CASCADE,
    -- Original natural language query text
    query_text TEXT NOT NULL,
    -- =====
    -- LLM PARSING RESULTS (for accuracy tracking)
    -- =====

    -- Parsed intent structure extracted by LLM before Pinecone query
    -- Structure: { "outdoor": true, "messy": false, "competitive": false,
    --             "age_appropriate_for": "child_uuid", "energy_level": "high" }
    parsed_intent JSONB NOT NULL,
    -- Confidence score of LLM parsing (0-1)
    parse_confidence DECIMAL(4, 3),
    -- Whether clarification was requested from user
    clarification_requested BOOLEAN DEFAULT false,
    clarification_question TEXT,
    clarification_response TEXT,
    -- =====
    -- PINECONE QUERY DETAILS
    -- =====
```

```
-- Pinecone query ID for debugging/tracing
pinecone_query_id VARCHAR(255),  
  
-- Filters applied to Pinecone query (derived from parsed_intent)
-- Structure: { "category": {"$in": ["stem", "building"]}, "min_age": {"$lte": 7} }
pinecone_filters JSONB,  
  
-- Number of results requested from Pinecone
pinecone_top_k INTEGER DEFAULT 10,  
  
-- ======  
-- SEARCH RESULTS (IDs only - actual data in PostgreSQL)
-- ======  
  
-- Activity IDs returned (ordered by relevance)
result_activity_ids JSONB, -- Structure: ["uuid1", "uuid2", ...]  
  
-- Pinecone similarity scores for each result
result_scores JSONB, -- Structure: [0.95, 0.88, 0.82, ...]  
  
-- Number of results returned
result_count INTEGER,  
  
-- ======  
-- USER INTERACTION TRACKING (for search quality metrics)
-- ======  
  
-- Did user interact with results?
had_interaction BOOLEAN DEFAULT false,
clicked_activity_id UUID REFERENCES activities(id),
resulted_in_booking BOOLEAN DEFAULT false,  
  
-- User feedback on search quality (if provided)
feedback_helpful BOOLEAN,
feedback_comment TEXT,  
  
-- Did user fall back to structured filters?
fallback_to_filters BOOLEAN DEFAULT false,  
  
-- ======  
-- PERFORMANCE & DEBUGGING
-- ======  
  
-- LangSmith trace ID for AI decision auditing
langsmith_trace_id VARCHAR(255),  
  
-- Response time breakdown (milliseconds)
```

```

llm_parse_duration_ms INTEGER,      -- Time for LLM to parse intent
pinecone_query_duration_ms INTEGER, -- Time for Pinecone vector search
postgres_enrich_duration_ms INTEGER,-- Time to fetch full activity data from PG
total_duration_ms INTEGER,         -- Total end-to-end time

-- Error tracking
had_error BOOLEAN DEFAULT false,
error_type VARCHAR(50), -- 'llm_parse_error', 'pinecone_timeout', 'no_results'
error_message TEXT,

-- Audit timestamps
created_at TIMESTAMP WITH TIME ZONE DEFAULT CURRENT_TIMESTAMP
);

-- Index on user_id for user search history
CREATE INDEX idx_nl_search_user_id ON nl_search_logs(user_id);

-- Index on created_at for date range queries
CREATE INDEX idx_nl_search_created_at ON nl_search_logs(created_at);

-- GIN index on parsed_intent for analyzing common intents
CREATE INDEX idx_nl_search_intent ON nl_search_logs USING GIN (parsed_intent);

-- Index for analyzing conversion rates
CREATE INDEX idx_nl_search_conversion ON nl_search_logs(resulted_in_booking) WHERE resulted_in_booking = true;

-- Index for error analysis and debugging
CREATE INDEX idx_nl_search_errors ON nl_search_logs(had_error, error_type) WHERE had_error = true;

-- Index for search quality analysis (fallback rate)
CREATE INDEX idx_nl_search_fallback ON nl_search_logs(fallback_to_filters) WHERE fallback_to_filters = true;

-- Full-text search on query_text for pattern analysis
CREATE INDEX idx_nl_search_query ON nl_search_logs USING GIN (to_tsvector('english', query_text));

```

COMMENT ON TABLE nl_search_logs IS 'Analytics log for NL searches. Actual search performed in Pinecone.';

```

-- =====
-- PINECONE SYNC STATUS TABLE
-- Operational tracking for PostgreSQL ↔ Pinecone synchronization
-- Used by N8N workflows and monitoring dashboards
-- =====

CREATE TABLE pinecone_sync_status (
    -- Primary identifier
    id UUID PRIMARY KEY DEFAULT gen_random_uuid(),

    -- Entity type being synced
    entity_type VARCHAR(50) NOT NULL CHECK (entity_type IN ('activity', 'child', 'venue')),

    -- Reference to the entity (polymorphic)
    entity_id UUID NOT NULL,

```

```
-- Pinecone namespace where vector is stored
pinecone_namespace VARCHAR(100) NOT NULL,

-- Pinecone vector ID
pinecone_vector_id VARCHAR(255) NOT NULL,

-- Sync status
sync_status VARCHAR(20) DEFAULT 'pending'
    CHECK (sync_status IN ('pending', 'synced', 'failed', 'deleted')),

-- Last successful sync
last_synced_at TIMESTAMP WITH TIME ZONE,

-- Last sync attempt (even if failed)
last_attempt_at TIMESTAMP WITH TIME ZONE,

-- Failure tracking
failure_count INTEGER DEFAULT 0,
last_error_message TEXT,

-- Hash of source data (to detect changes without re-embedding)
content_hash VARCHAR(64),

-- Embedding model version
embedding_model_version VARCHAR(50),

-- Vector dimensions (for validation)
vector_dimensions INTEGER,

-- Audit timestamps
created_at TIMESTAMP WITH TIME ZONE DEFAULT CURRENT_TIMESTAMP,
updated_at TIMESTAMP WITH TIME ZONE DEFAULT CURRENT_TIMESTAMP,

-- Ensure one sync record per entity
CONSTRAINT unique_sync_entity UNIQUE (entity_type, entity_id)
);

-- Index for finding entities needing sync
CREATE INDEX idx_pinecone_sync_pending ON pinecone_sync_status(sync_status, entity_type)
    WHERE sync_status IN ('pending', 'failed');

-- Index for monitoring failed syncs
CREATE INDEX idx_pinecone_sync_failed ON pinecone_sync_status(failure_count)
    WHERE sync_status = 'failed';

-- Index on namespace for bulk operations
CREATE INDEX idx_pinecone_sync_namespace ON pinecone_sync_status(pinecone_namespace);
```

```
COMMENT ON TABLE pinecone_sync_status IS 'Tracks synchronization status between PostgreSQL and Pinecone vectors'
```

2.10 Audit and Compliance Tables

sql

```
-- =====
-- AUDIT LOGS TABLE
-- Comprehensive audit trail for compliance (COPPA/GDPR-K)
-- Retains logs for 2 years per PRD Section 10
-- =====

CREATE TABLE audit_logs (
    -- Primary identifier
    id UUID PRIMARY KEY DEFAULT gen_random_uuid(),

    -- User who performed the action (NULL for system actions)
    user_id UUID REFERENCES users(id) ON DELETE SET NULL,

    -- Action category
    action_category VARCHAR(50) NOT NULL,
    -- Categories: 'auth', 'profile', 'booking', 'child', 'venue', 'admin', 'ai', 'system'

    -- Specific action performed
    action_type VARCHAR(100) NOT NULL,
    -- Examples: 'login', 'logout', 'profile_update', 'child_create', 'booking_create',
    --           'booking_cancel', 'consent_given', 'data_export', 'data_delete'

    -- Target entity type (polymorphic)
    entity_type VARCHAR(50),
    entity_id UUID,

    -- Previous and new values for change tracking (sensitive data encrypted)
    old_values JSONB,
    new_values JSONB,

    -- Request metadata for security analysis
    ip_address INET,
    user_agent TEXT,

    -- Session ID for grouping related actions
    session_id VARCHAR(255),

    -- Geographic location (derived from IP if available)
    geo_country VARCHAR(100),
    geo_city VARCHAR(100),

    -- Request ID for correlating with application logs
    request_id VARCHAR(255),

    -- Severity level
    severity VARCHAR(20) DEFAULT 'info' CHECK (severity IN ('debug', 'info', 'warning', 'error', 'critical')),
```

```
-- Additional context as flexible JSONB
context JSONB DEFAULT '{}':jsonb,
-- Timestamp (immutable)
created_at TIMESTAMP WITH TIME ZONE DEFAULT CURRENT_TIMESTAMP
);

-- Index on user_id for user activity queries
CREATE INDEX idx_audit_logs_user_id ON audit_logs(user_id);
-- Index on action_category for filtering
CREATE INDEX idx_audit_logs_category ON audit_logs(action_category);
-- Index on action_type for specific action queries
CREATE INDEX idx_audit_logs_action ON audit_logs(action_type);
-- Index on entity for finding all actions on an entity
CREATE INDEX idx_audit_logs_entity ON audit_logs(entity_type, entity_id);
-- Index on created_at for date range queries and retention
CREATE INDEX idx_audit_logs_created_at ON audit_logs(created_at);
-- Index on severity for error analysis
CREATE INDEX idx_audit_logs_severity ON audit_logs(severity) WHERE severity IN ('warning', 'error', 'critical');
-- Composite index for compliance queries (user activity over time)
CREATE INDEX idx_audit_logs_compliance ON audit_logs(user_id, created_at DESC);
```

```
COMMENT ON TABLE audit_logs IS 'Comprehensive audit trail for COPPA/GDPR-K compliance (2-year retention);'
```

```
-- =====
-- PARENTAL CONSENT RECORDS TABLE
-- Explicit tracking of COPPA-required parental consent for child data
-- =====
CREATE TABLE parental_consent_records (
    -- Primary identifier
    id UUID PRIMARY KEY DEFAULT gen_random_uuid(),
    -- Reference to the parent giving consent
    parent_id UUID NOT NULL REFERENCES users(id) ON DELETE CASCADE,
    -- Reference to the child consent is for
    child_id UUID NOT NULL REFERENCES children(id) ON DELETE CASCADE,
    -- Type of consent
    consent_type VARCHAR(50) NOT NULL,
    -- Types: 'data_collection', 'data_processing', 'activity_booking', 'photo_sharing'
    -- Consent version (tracks consent form version)
    consent_version VARCHAR(20) NOT NULL,
    -- Was consent given?
    consent_given BOOLEAN NOT NULL DEFAULT false
);
```

```

consent_given BOOLEAN NOT NULL,
-- Consent method
consent_method VARCHAR(50) NOT NULL, -- 'checkbox', 'signature', 'email_verification'

-- Full consent text that was agreed to
consent_text TEXT NOT NULL,

-- IP address at time of consent
ip_address INET,

-- User agent at time of consent
user_agent TEXT,

-- Withdrawal tracking
withdrawn BOOLEAN DEFAULT false,
withdrawn_at TIMESTAMP WITH TIME ZONE,
withdrawal_reason TEXT,

-- Timestamps
consented_at TIMESTAMP WITH TIME ZONE DEFAULT CURRENT_TIMESTAMP,

-- Ensure one consent record per type per child (latest version)
CONSTRAINT unique_consenter_record UNIQUE (child_id, consent_type, consent_version)
);

-- Index on parent_id for consent management
CREATE INDEX idx_consenter_parent_id ON parental_consenter_records(parent_id);
-- Index on child_id for compliance checks
CREATE INDEX idx_consenter_child_id ON parental_consenter_records(child_id);
-- Index for active consents
CREATE INDEX idx_consenter_active ON parental_consenter_records(child_id, consent_type)
    WHERE consent_given = true AND withdrawn = false;

COMMENT ON TABLE parental_consenter_records IS 'COPPA verifiable parental consent tracking';

```

2.11 Platform Configuration Tables

sql

```

-- =====
-- PLATFORM SETTINGS TABLE
-- System-wide configuration parameters managed by platform admins
-- =====

CREATE TABLE platform_settings (
    -- Setting key (unique identifier)
    key VARCHAR(100) PRIMARY KEY,
    -- Setting value (stored as JSONB for flexibility)
    value JSONB NOT NULL,
    -- Human-readable description
    description TEXT,
    -- Setting category for grouping
    category VARCHAR(50),
    -- Data type hint for UI
    data_type VARCHAR(20) CHECK (data_type IN ('string', 'number', 'boolean', 'json', 'array')),
    -- Is this setting publicly visible?
    is_public BOOLEAN DEFAULT false,
    -- Last modified tracking
    updated_by UUID REFERENCES users(id),
    updated_at TIMESTAMP WITH TIME ZONE DEFAULT CURRENT_TIMESTAMP
);

-- Seed essential platform settings
INSERT INTO platform_settings (key, value, description, category, data_type) VALUES
('platform_fee_default', '17.5', 'Default platform fee percentage for new venues', 'financial', 'number'),
('cancellation_window_hours', '48', 'Hours before activity for full refund eligibility', 'booking', 'number'),
('waitlist_confirmation_hours', '4', 'Hours given to confirm waitlist spot', 'booking', 'number'),
('credit_expiry_warning_days', '7', 'Days before month end to send expiry warning', 'credits', 'number'),
('venue_score_curation_threshold', '60', 'Minimum venue score for AI curation eligibility', 'curation', 'number'),
('maxSuggestionsPerMonth', '3', 'Maximum AI suggestions per child per month', 'curation', 'number'),
('feedbackRequestDelayHours', '1', 'Hours after activity to send feedback request', 'feedback', 'number'),
('dataRetentionDays', '730', 'Days to retain audit logs (2 years)', 'compliance', 'number'),
('childDeletionDelayDays', '30', 'Days after deletion request to purge child data', 'compliance', 'number');

COMMENT ON TABLE platform_settings IS 'System-wide configuration parameters';

-- =====
-- FEATURE FLAGS TABLE
-- Runtime feature toggles for gradual rollouts and A/B testing (UI Spec line 803)
-- =====

```

```
-- =====
CREATE TABLE feature_flags (
    -- Primary identifier
    id UUID PRIMARY KEY DEFAULT gen_random_uuid(),
    -- Feature flag key (unique identifier used in code)
    flag_key VARCHAR(100) NOT NULL UNIQUE,
    -- Human-readable name and description
    name VARCHAR(255) NOT NULL,
    description TEXT,
    -- Feature category for grouping
    category VARCHAR(50), -- 'ai', 'booking', 'payment', 'ui', 'integration'
    -- Global enable/disable
    is_enabled BOOLEAN DEFAULT false,
    -- Percentage rollout (0-100, for gradual rollouts)
    rollout_percentage INTEGER DEFAULT 0 CHECK (rollout_percentage BETWEEN 0 AND 100),
    -- Target specific user roles (NULL means all roles)
    -- Structure: ["parent", "hr_admin", "venue_admin"]
    target_roles JSONB,
    -- Target specific corporate accounts (for beta testing with specific companies)
    target_corporate_ids JSONB DEFAULT '[]'::jsonb,
    -- Target specific users (for internal testing)
    target_user_ids JSONB DEFAULT '[]'::jsonb,
    -- Environment targeting
    environments JSONB DEFAULT '["production", "staging", "development"]'::jsonb,
    -- Feature flag variants for A/B testing
    -- Structure: { "control": 50, "variant_a": 25, "variant_b": 25 }
    variants JSONB,
    -- Start and end dates for time-limited features
    starts_at TIMESTAMP WITH TIME ZONE,
    ends_at TIMESTAMP WITH TIME ZONE,
    -- Audit tracking
    created_by UUID REFERENCES users(id),
    updated_by UUID REFERENCES users(id),
    created_at TIMESTAMP WITH TIME ZONE DEFAULT CURRENT_TIMESTAMP,
    updated_at TIMESTAMP WITH TIME ZONE DEFAULT CURRENT_TIMESTAMP
```

```

);

-- Index on flag_key for fast lookups
CREATE INDEX idx_feature_flags_key ON feature_flags(flag_key);
-- Index on enabled flags for runtime queries
CREATE INDEX idx_feature_flags_enabled ON feature_flags(is_enabled) WHERE is_enabled = true;
-- Index on category for admin filtering
CREATE INDEX idx_feature_flags_category ON feature_flags(category);

COMMENT ON TABLE feature_flags IS 'Runtime feature toggles for gradual rollouts and A/B testing';

```

```

-- =====
-- NOTIFICATION TEMPLATES TABLE
-- Configurable notification content templates (UI Spec line 802)
-- =====

CREATE TABLE notification_templates (
    -- Primary identifier
    id UUID PRIMARY KEY DEFAULT gen_random_uuid(),

    -- Template identifier used in code
    template_key VARCHAR(100) NOT NULL,

    -- Notification channel this template is for
    channel notification_channel NOT NULL,

    -- Language code for internationalization
    language_code VARCHAR(10) DEFAULT 'en',

    -- Template name for admin reference
    name VARCHAR(255) NOT NULL,

    -- Email-specific fields
    email_subject VARCHAR(255),
    email_body_html TEXT,
    email_body_text TEXT,

    -- SMS/WhatsApp/Push body
    short_body TEXT,

    -- In-app notification
    in_app_title VARCHAR(255),
    in_app_body TEXT,

    -- Available merge variables for this template
    -- Structure: ["parent_name", "child_name", "activity_name", "booking_date"]
    available_variables JSONB DEFAULT '[]'::jsonb,

```

```

-- Template category for grouping
category VARCHAR(50), -- 'booking', 'reminder', 'waitlist', 'credits', 'system'

-- Is this the active version?
is_active BOOLEAN DEFAULT true,

-- Version tracking for template history
version INTEGER DEFAULT 1,

-- Audit timestamps
created_at TIMESTAMP WITH TIME ZONE DEFAULT CURRENT_TIMESTAMP,
updated_at TIMESTAMP WITH TIME ZONE DEFAULT CURRENT_TIMESTAMP,

-- Ensure unique template per key, channel, and language
CONSTRAINT unique_template_key_channel_lang UNIQUE (template_key, channel, language_code)
);

-- Index on template_key for lookups
CREATE INDEX idx_notification_templates_key ON notification_templates(template_key);

-- Index on channel for filtering
CREATE INDEX idx_notification_templates_channel ON notification_templates(channel);

-- Index on active templates
CREATE INDEX idx_notification_templates_active ON notification_templates(is_active) WHERE is_active = true;

COMMENT ON TABLE notification_templates IS 'Configurable notification content templates with multi-language support';

=====
-- DATA EXPORT REQUESTS TABLE
-- GDPR/CCPA data export tracking (UI Spec line 414 - "Download my data")
=====

CREATE TABLE data_export_requests (
    -- Primary identifier
    id UUID PRIMARY KEY DEFAULT gen_random_uuid(),

    -- Reference to the user requesting export
    user_id UUID NOT NULL REFERENCES users(id) ON DELETE CASCADE,

    -- Type of export
    export_type VARCHAR(50) NOT NULL, -- 'full_data', 'children_data', 'booking_history', 'personal_info'

    -- Export format
    format VARCHAR(20) DEFAULT 'json', -- 'json', 'csv', 'pdf'

    -- Request status
    status VARCHAR(20) DEFAULT 'pending'
)

```

```

CHECK (status IN ('pending', 'processing', 'completed', 'failed', 'expired')),

-- Processing timestamps
requested_at TIMESTAMP WITH TIME ZONE DEFAULT CURRENT_TIMESTAMP,
started_at TIMESTAMP WITH TIME ZONE,
completed_at TIMESTAMP WITH TIME ZONE,

-- Generated file details
file_url VARCHAR(500),
file_size_bytes BIGINT,
file_expires_at TIMESTAMP WITH TIME ZONE, -- Auto-delete after X days

-- Download tracking
download_count INTEGER DEFAULT 0,
last_downloaded_at TIMESTAMP WITH TIME ZONE,

-- Error tracking
error_message TEXT,

-- Request metadata (IP, user agent for compliance)
request_ip INET,
request_user_agent TEXT,

-- Audit timestamps
created_at TIMESTAMP WITH TIME ZONE DEFAULT CURRENT_TIMESTAMP
);

-- Index on user_id for listing user's export requests
CREATE INDEX idx_data_exports_user_id ON data_export_requests(user_id);
-- Index on status for processing queue
CREATE INDEX idx_data_exports_status ON data_export_requests(status) WHERE status IN ('pending', 'processing');
-- Index on expiry for cleanup job
CREATE INDEX idx_data_exports_expiry ON data_export_requests(file_expires_at)
    WHERE status = 'completed' AND file_expires_at IS NOT NULL;

COMMENT ON TABLE data_export_requests IS 'GDPR/CCPA data export request tracking';

=====

-- CHAT CONVERSATIONS TABLE
-- Stores NL search conversation sessions for context and history
-- Supports multi-turn conversations with the AI search interface
=====

CREATE TABLE chat_conversations (
    -- Primary identifier
    id UUID PRIMARY KEY DEFAULT gen_random_uuid(),

```

```

-- Reference to the user
user_id UUID NOT NULL REFERENCES users(id) ON DELETE CASCADE,

-- Conversation title (auto-generated or user-set)
title VARCHAR(255),

-- Conversation type
conversation_type VARCHAR(50) DEFAULT 'activity_search', -- 'activity_search', 'support', 'booking_help'

-- Child context (if searching for specific child)
child_id UUID REFERENCES children(id) ON DELETE SET NULL,

-- Is conversation still active?
is_active BOOLEAN DEFAULT true,

-- Timestamps
started_at TIMESTAMP WITH TIME ZONE DEFAULT CURRENT_TIMESTAMP,
last_message_at TIMESTAMP WITH TIME ZONE,
ended_at TIMESTAMP WITH TIME ZONE,

-- Conversation summary (AI-generated for quick reference)
summary TEXT,

-- Outcome tracking
resulted_in_booking BOOLEAN DEFAULT false,
booking_id UUID REFERENCES bookings(id) ON DELETE SET NULL,

-- Audit timestamps
created_at TIMESTAMP WITH TIME ZONE DEFAULT CURRENT_TIMESTAMP,
updated_at TIMESTAMP WITH TIME ZONE DEFAULT CURRENT_TIMESTAMP
);

-- Index on user_id for user's conversation history
CREATE INDEX idx_chat_conversations_user_id ON chat_conversations(user_id);
-- Index on active conversations
CREATE INDEX idx_chat_conversations_active ON chat_conversations(user_id, is_active, last_message_at DESC)
    WHERE is_active = true;
-- Index on child_id for child-specific conversations
CREATE INDEX idx_chat_conversations_child ON chat_conversations(child_id) WHERE child_id IS NOT NULL;

COMMENT ON TABLE chat_conversations IS 'NL search conversation sessions for multi-turn AI interactions';

-- =====
-- CHAT MESSAGES TABLE
-- Individual messages within a conversation
-- =====

```

```

CREATE TABLE chat_messages (
    -- Primary identifier
    id UUID PRIMARY KEY DEFAULT gen_random_uuid(),

    -- Reference to the conversation
    conversation_id UUID NOT NULL REFERENCES chat_conversations(id) ON DELETE CASCADE,

    -- Message role
    role VARCHAR(20) NOT NULL CHECK (role IN ('user', 'assistant', 'system')),

    -- Message content
    content TEXT NOT NULL,

    -- For assistant messages: parsed intent (if applicable)
    parsed_intent JSONB,

    -- Activity results shown (if search was performed)
    result_activity_ids JSONB,

    -- Quick action chips shown to user
    suggested_actions JSONB,

    -- User interaction with this message
    user_clicked_activity_id UUID REFERENCES activities(id),
    user_selected_action VARCHAR(100),

    -- LangSmith trace for AI messages
    langsmith_trace_id VARCHAR(255),

    -- Timing
    created_at TIMESTAMP WITH TIME ZONE DEFAULT CURRENT_TIMESTAMP,

    -- Message ordering within conversation
    sequence_number INTEGER NOT NULL
);

-- Index on conversation_id for fetching messages
CREATE INDEX idx_chat_messages_conversation ON chat_messages(conversation_id, sequence_number);

-- Index on role for filtering
CREATE INDEX idx_chat_messages_role ON chat_messages(conversation_id, role);

COMMENT ON TABLE chat_messages IS 'Individual messages in NL search conversations';

-- =====
-- API KEYS TABLE
-- API key management for platform admin (UI Spec line 804)

```

```
-- =====
CREATE TABLE api_keys (
    -- Primary identifier
    id UUID PRIMARY KEY DEFAULT gen_random_uuid(),
    -- Human-readable name for the API key
    name VARCHAR(255) NOT NULL,
    -- Description of what this key is used for
    description TEXT,
    -- The actual API key (hashed, only shown once on creation)
    key_hash VARCHAR(255) NOT NULL,
    -- Key prefix for identification (first 8 chars, stored unhashed)
    key_prefix VARCHAR(10) NOT NULL,
    -- Key type/purpose
    key_type VARCHAR(50) NOT NULL, -- 'integration', 'testing', 'service', 'webhook'
    -- Permissions/scopes granted to this key
    -- Structure: ["read:activities", "write:bookings", "admin:venues"]
    scopes JSONB DEFAULT '[]'::jsonb,
    -- Rate limiting configuration
    rate_limit_per_minute INTEGER DEFAULT 60,
    rate_limit_per_day INTEGER DEFAULT 10000,
    -- IP whitelist (NULL means no restriction)
    allowed_ips JSONB, -- Structure: ["192.168.1.1", "10.0.0.0/8"]
    -- Expiration
    expires_at TIMESTAMP WITH TIME ZONE,
    -- Status
    is_active BOOLEAN DEFAULT true,
    revoked_at TIMESTAMP WITH TIME ZONE,
    revoked_by UUID REFERENCES users(id),
    revoke_reason TEXT,
    -- Usage tracking
    last_used_at TIMESTAMP WITH TIME ZONE,
    last_used_ip INET,
    total_requests BIGINT DEFAULT 0,
    -- Created by (platform admin)
    created_by UUID NOT NULL REFERENCES users(id),
```

```

-- Audit timestamps
created_at TIMESTAMP WITH TIME ZONE DEFAULT CURRENT_TIMESTAMP,
updated_at TIMESTAMP WITH TIME ZONE DEFAULT CURRENT_TIMESTAMP
);

-- Index on key_prefix for lookup (full hash comparison done after)
CREATE INDEX idx_api_keys_prefix ON api_keys(key_prefix) WHERE is_active = true;
-- Index on key_type for filtering
CREATE INDEX idx_api_keys_type ON api_keys(key_type);
-- Index on expiration for cleanup
CREATE INDEX idx_api_keys_expiry ON api_keys(expires_at) WHERE is_active = true AND expires_at IS NOT NULL;

COMMENT ON TABLE api_keys IS 'API key management for integrations and external access';

```

```

-- =====
-- INVOICES TABLE
-- Corporate billing invoices generated monthly
-- =====

CREATE TABLE invoices (
    -- Primary identifier
    id UUID PRIMARY KEY DEFAULT gen_random_uuid(),

    -- Invoice number for display (sequential per corporate)
    invoice_number VARCHAR(50) NOT NULL UNIQUE,

    -- Reference to the corporate account
    corporate_id UUID NOT NULL REFERENCES corporates(id) ON DELETE RESTRICT,

    -- Reference to the subscription
    subscription_id UUID NOT NULL REFERENCES corporate_subscriptions(id) ON DELETE RESTRICT,

    -- Billing period
    billing_period_start DATE NOT NULL,
    billing_period_end DATE NOT NULL,

    -- Invoice amounts
    subtotal DECIMAL(12, 2) NOT NULL CHECK (subtotal >= 0),
    tax_amount DECIMAL(12, 2) DEFAULT 0 CHECK (tax_amount >= 0),
    total_amount DECIMAL(12, 2) NOT NULL CHECK (total_amount >= 0),

    -- Currency
    currency VARCHAR(3) DEFAULT 'USD',

    -- Invoice status
    status VARCHAR(20) DEFAULT 'draft' CHECK (status IN ('draft', 'sent', 'paid', 'overdue', 'cancelled')),

```

```

-- Due date (net-30 by default)
due_date DATE NOT NULL,

-- Payment tracking
paid_at TIMESTAMP WITH TIME ZONE,
paid_amount DECIMAL(12, 2),
payment_method VARCHAR(50),
payment_reference VARCHAR(255),

-- Stripe invoice ID
stripe_invoice_id VARCHAR(255),

-- Invoice PDF URL
pdf_url VARCHAR(500),

-- Line items stored as JSONB for flexibility
line_items JSONB NOT NULL,

-- Notes
notes TEXT,

-- Sent tracking
sent_at TIMESTAMP WITH TIME ZONE,

-- Audit timestamps
created_at TIMESTAMP WITH TIME ZONE DEFAULT CURRENT_TIMESTAMP,
updated_at TIMESTAMP WITH TIME ZONE DEFAULT CURRENT_TIMESTAMP
);

-- Index on corporate_id for retrieving invoices
CREATE INDEX idx_invoices_corporate_id ON invoices(corporate_id);

-- Index on status for filtering
CREATE INDEX idx_invoices_status ON invoices(status);

-- Index on due_date for overdue processing
CREATE INDEX idx_invoices_due_date ON invoices(due_date) WHERE status = 'sent';

COMMENT ON TABLE invoices IS 'Corporate billing invoices with Stripe integration';

```

2.12 Update Trigger Function

sql

```
-- =====
-- AUTOMATIC UPDATED_AT TRIGGER FUNCTION
-- Automatically updates the updated_at timestamp on row modification
-- =====

CREATE OR REPLACE FUNCTION update_updated_at_column()
RETURNS TRIGGER AS $$

BEGIN
    NEW.updated_at = CURRENT_TIMESTAMP;
    RETURN NEW;
END;

$$ language 'plpgsql';

-- Apply trigger to all tables with updated_at column
CREATE TRIGGER update_corporates_updated_at BEFORE UPDATE ON corporates
    FOR EACH ROW EXECUTE FUNCTION update_updated_at_column();

CREATE TRIGGER update_users_updated_at BEFORE UPDATE ON users
    FOR EACH ROW EXECUTE FUNCTION update_updated_at_column();

CREATE TRIGGER update_user_locations_updated_at BEFORE UPDATE ON user_locations
    FOR EACH ROW EXECUTE FUNCTION update_updated_at_column();

CREATE TRIGGER update_venues_updated_at BEFORE UPDATE ON venues
    FOR EACH ROW EXECUTE FUNCTION update_updated_at_column();

CREATE TRIGGER update_activities_updated_at BEFORE UPDATE ON activities
    FOR EACH ROW EXECUTE FUNCTION update_updated_at_column();

CREATE TRIGGER update_activity_sessions_updated_at BEFORE UPDATE ON activity_sessions
    FOR EACH ROW EXECUTE FUNCTION update_updated_at_column();

CREATE TRIGGER update_children_updated_at BEFORE UPDATE ON children
    FOR EACH ROW EXECUTE FUNCTION update_updated_at_column();

CREATE TRIGGER update_corporate_subscriptions_updated_at BEFORE UPDATE ON corporate_subscriptions
    FOR EACH ROW EXECUTE FUNCTION update_updated_at_column();

CREATE TRIGGER update_user_credit_balances_updated_at BEFORE UPDATE ON user_credit_balances
    FOR EACH ROW EXECUTE FUNCTION update_updated_at_column();

CREATE TRIGGER update_bookings_updated_at BEFORE UPDATE ON bookings
    FOR EACH ROW EXECUTE FUNCTION update_updated_at_column();

CREATE TRIGGER update_waitlist_updated_at BEFORE UPDATE ON waitlist
    FOR EACH ROW EXECUTE FUNCTION update_updated_at_column();
```

```
CREATE TRIGGER update_reviews_updated_at BEFORE UPDATE ON reviews
  FOR EACH ROW EXECUTE FUNCTION update_updated_at_column();
```

```
CREATE TRIGGER update_ai_suggestions_updated_at BEFORE UPDATE ON ai_suggestions
  FOR EACH ROW EXECUTE FUNCTION update_updated_at_column();
```

```
CREATE TRIGGER update_invoices_updated_at BEFORE UPDATE ON invoices
  FOR EACH ROW EXECUTE FUNCTION update_updated_at_column();
```

```
CREATE TRIGGER update_pinecone_sync_status_updated_at BEFORE UPDATE ON pinecone_sync_status
  FOR EACH ROW EXECUTE FUNCTION update_updated_at_column();
```

```
CREATE TRIGGER update_feature_flags_updated_at BEFORE UPDATE ON feature_flags
  FOR EACH ROW EXECUTE FUNCTION update_updated_at_column();
```

```
CREATE TRIGGER update_notification_templates_updated_at BEFORE UPDATE ON notification_templates
  FOR EACH ROW EXECUTE FUNCTION update_updated_at_column();
```

```
CREATE TRIGGER update_chat_conversations_updated_at BEFORE UPDATE ON chat_conversations
  FOR EACH ROW EXECUTE FUNCTION update_updated_at_column();
```

```
CREATE TRIGGER update_api_keys_updated_at BEFORE UPDATE ON api_keys
  FOR EACH ROW EXECUTE FUNCTION update_updated_at_column();
```

```
CREATE TRIGGER update_employee_invitations_updated_at BEFORE UPDATE ON employee_invitations
  FOR EACH ROW EXECUTE FUNCTION update_updated_at_column();
```

```
CREATE TRIGGER update_activity_recurring_schedules_updated_at BEFORE UPDATE ON activity_recurring_schedules
  FOR EACH ROW EXECUTE FUNCTION update_updated_at_column();
```

```
CREATE TRIGGER update_venue_payouts_updated_at BEFORE UPDATE ON venue_payouts
  FOR EACH ROW EXECUTE FUNCTION update_updated_at_column();
```

```
CREATE TRIGGER update_venue_integrations_updated_at BEFORE UPDATE ON venue_integrations
  FOR EACH ROW EXECUTE FUNCTION update_updated_at_column();
```

```
-- =====
```

```
-- TRIGGER: Mark activities for Pinecone re-sync on content changes
```

```
-- =====
```

```
CREATE OR REPLACE FUNCTION mark_activity_for_pinecone_sync()
```

```
RETURNS TRIGGER AS $$
```

```
BEGIN
```

```
-- Only mark for sync if searchable content changed
```

```
IF (OLD.name IS DISTINCT FROM NEW.name OR
```

```
  OLD.short_description IS DISTINCT FROM NEW.short_description OR
```

```
  OLD.full_description IS DISTINCT FROM NEW.full_description OR
```

```

OLD.category IS DISTINCT FROM NEW.category OR
OLD.activity_tags IS DISTINCT FROM NEW.activity_tags) THEN
    NEW.pinecone_sync_required = true;
END IF;
RETURN NEW;
END;
$$ LANGUAGE plpgsql;

```

```

CREATE TRIGGER trigger_activity_pinecone_sync BEFORE UPDATE ON activities
FOR EACH ROW EXECUTE FUNCTION mark_activity_for_pinecone_sync();

```

```

-- =====
-- TRIGGER: Mark children for Pinecone re-sync on preference changes
-- =====

CREATE OR REPLACE FUNCTION mark_child_for_pinecone_sync()
RETURNS TRIGGER AS $$

BEGIN
    -- Only mark for sync if preference-related fields changed
    IF (OLD.energy_level IS DISTINCT FROM NEW.energy_level OR
        OLD.social_preference IS DISTINCT FROM NEW.social_preference OR
        OLD.interests IS DISTINCT FROM NEW.interests OR
        OLD.activities_to_avoid IS DISTINCT FROM NEW.activities_to_avoid) THEN
        NEW.pinecone_sync_required = true;
    END IF;
    RETURN NEW;
END;
$$ LANGUAGE plpgsql;

```

```

CREATE TRIGGER trigger_child_pinecone_sync BEFORE UPDATE ON children
FOR EACH ROW EXECUTE FUNCTION mark_child_for_pinecone_sync();

```

3. Performance Optimizations

3.1 Materialized View for Dashboard Queries

sql

```

-- =====
-- MATERIALIZED VIEW: PARENT DASHBOARD SUMMARY
-- Pre-computed view for fast parent dashboard loading
-- Refresh hourly via N8N workflow
-- =====

CREATE MATERIALIZED VIEW parent_dashboard_summary AS
SELECT
    u.id AS user_id,
    u.first_name,
    u.last_name,
    ucb.credits_available,
    ucb.credits_used,
    ucb.credits_allocated,
    ucb.billing_period_start,
    ucb.billing_period_end,
    -- Days remaining calculation
    GREATEST(0, ucb.billing_period_end - CURRENT_DATE) AS days_remaining,
    -- Count of children
    (SELECT COUNT(*) FROM children c WHERE c.parent_id = u.id AND c.is_active = true) AS child_count,
    -- Count of upcoming bookings
    (SELECT COUNT(*)
     FROM bookings b
     JOIN activity_sessions s ON b.session_id = s.id
     WHERE b.booked_by = u.id
       AND b.status IN ('pending', 'confirmed')
       AND s.session_date >= CURRENT_DATE) AS upcoming_booking_count,
    -- Count of pending suggestions
    (SELECT COUNT(*)
     FROM ai_suggestions ais
     WHERE ais.user_id = u.id
       AND ais.parent_response IS NULL
       AND ais.expires_at > CURRENT_TIMESTAMP) AS pending_suggestion_count,
    -- Count of unread notifications
    (SELECT COUNT(*)
     FROM notifications n
     WHERE n.user_id = u.id
       AND n.is_read = false) AS unread_notification_count
FROM users u
LEFT JOIN user_credit_balances ucb ON ucb.user_id = u.id
  AND ucb.billing_period_start <= CURRENT_DATE
  AND ucb.billing_period_end >= CURRENT_DATE
WHERE u.role = 'parent' AND u.is_active = true;

-- Index on the materialized view
CREATE UNIQUE INDEX idx_dashboard_summary_user_id ON parent_dashboard_summary(user_id);

```

```
-- Function to refresh the materialized view
CREATE OR REPLACE FUNCTION refresh_parent_dashboard_summary()
RETURNS void AS $$

BEGIN
    REFRESH MATERIALIZED VIEW CONCURRENTLY parent_dashboard_summary;
END;

$$ LANGUAGE plpgsql;

COMMENT ON MATERIALIZED VIEW parent_dashboard_summary IS 'Pre-computed parent dashboard data refreshed hourly';
```

3.2 Partial Indexes for Common Query Patterns

```
sql

-- =====
-- ADDITIONAL PARTIAL INDEXES FOR PERFORMANCE
-- Optimized for the most common query patterns identified in UI Spec
-- =====

-- Fast lookup of active activities by category and age (Browse Activities screen)
CREATE INDEX idx_activities_browse ON activities(category, min_age, max_age, average_rating DESC)
    WHERE is_active = true AND is_published = true;

-- Fast lookup of available sessions in the next 30 days
CREATE INDEX idx_sessions_next_30_days ON activity_sessions(session_date, activity_id)
    WHERE session_date BETWEEN CURRENT_DATE AND CURRENT_DATE + INTERVAL '30 days'
        AND is_cancelled = false;

-- Fast lookup of bookings needing reminder (24 hours before)
CREATE INDEX idx_bookingsReminder_24h ON bookings(id)
    WHERE status = 'confirmed' AND reminder_24h_sent = false;

-- Fast lookup of expired waitlist offers
CREATE INDEX idx_waitlist_expired ON waitlist(spot_expires_at)
    WHERE is_active = true AND spot_offered_at IS NOT NULL AND response IS NULL;
```

4. Database Maintenance

4.1 Data Retention Policies

```
sql
```

```

-- =====
-- DATA RETENTION AND CLEANUP
-- Implements retention policies per PRD compliance requirements
-- =====

-- Function to purge old audit logs (retain 2 years)
CREATE OR REPLACE FUNCTION purge_old_audit_logs()
RETURNS INTEGER AS $$

DECLARE
    deleted_count INTEGER;
BEGIN
    DELETE FROM audit_logs
    WHERE created_at < CURRENT_TIMESTAMP - INTERVAL '2 years';
    GET DIAGNOSTICS deleted_count = ROW_COUNT;
    RETURN deleted_count;
END;
$$ LANGUAGE plpgsql;

-- Function to purge deleted children data (30 days after request)
CREATE OR REPLACE FUNCTION purge_deleted_children()
RETURNS INTEGER AS $$

DECLARE
    deleted_count INTEGER;
BEGIN
    DELETE FROM children
    WHERE deletion_requested = true
    AND deletion_scheduled_for <= CURRENT_DATE;
    GET DIAGNOSTICS deleted_count = ROW_COUNT;
    RETURN deleted_count;
END;
$$ LANGUAGE plpgsql;

-- Function to purge old NL search logs (retain 90 days)
CREATE OR REPLACE FUNCTION purge_old_search_logs()
RETURNS INTEGER AS $$

DECLARE
    deleted_count INTEGER;
BEGIN
    DELETE FROM nl_search_logs
    WHERE created_at < CURRENT_TIMESTAMP - INTERVAL '90 days';
    GET DIAGNOSTICS deleted_count = ROW_COUNT;
    RETURN deleted_count;
END;
$$ LANGUAGE plpgsql;

```

COMMENT ON FUNCTION purge_old_audit_logs IS 'Removes audit logs older than 2 years';

```
COMMENT ON FUNCTION purge_deleted_children IS 'Purges child data 30 days after deletion request';
COMMENT ON FUNCTION purge_old_search_logs IS 'Removes NL search logs older than 90 days';
```

-- Function to purge expired data export files (retain 7 days)

```
CREATE OR REPLACE FUNCTION purge_expired_data_exports()
RETURNS INTEGER AS $$

DECLARE
    deleted_count INTEGER;

BEGIN
    UPDATE data_export_requests
    SET status = 'expired', file_url = NULL
    WHERE status = 'completed'
    AND file_expires_at < CURRENT_TIMESTAMP;
    GET DIAGNOSTICS deleted_count = ROW_COUNT;
    RETURN deleted_count;
END;
$$ LANGUAGE plpgsql;
```

-- Function to purge old chat conversations (retain 90 days)

```
CREATE OR REPLACE FUNCTION purge_old_chat_conversations()
RETURNS INTEGER AS $$

DECLARE
    deleted_count INTEGER;

BEGIN
    DELETE FROM chat_conversations
    WHERE is_active = false
    AND ended_at < CURRENT_TIMESTAMP - INTERVAL '90 days';
    GET DIAGNOSTICS deleted_count = ROW_COUNT;
    RETURN deleted_count;
END;
$$ LANGUAGE plpgsql;
```

-- Function to expire pending employee invitations

```
CREATE OR REPLACE FUNCTION expire_pending_invitations()
RETURNS INTEGER AS $$

DECLARE
    updated_count INTEGER;

BEGIN
    UPDATE employee_invitations
    SET status = 'expired'
    WHERE status = 'pending'
    AND expires_at < CURRENT_TIMESTAMP;
    GET DIAGNOSTICS updated_count = ROW_COUNT;
    RETURN updated_count;
END;
$$ LANGUAGE plpgsql;
```

COMMENT ON FUNCTION purge_expired_data_exports IS 'Marks expired data exports and removes file URLs';
COMMENT ON FUNCTION purge_old_chat_conversations IS 'Removes inactive chat conversations older than 90 days';
COMMENT ON FUNCTION expire_pending_invitations IS 'Expires pending employee invitations past their expiry date';

5. Summary Statistics

Metric	Count
Total Tables	32
Enum Types	13
Indexes	95+
Foreign Keys	45+
Triggers	27
Materialized Views	1
Utility Functions	9

Tables by Domain

Domain	Tables
Users & Auth	(users), (user_locations), (corporates), (employee_invitations)
Venues	(venues), (venue_performance_scores), (venue_payouts), (venue_integrations)
Activities	(activities), (activity_sessions), (activity_recurring_schedules)
Children	(children), (childFavorites)
Subscriptions	(corporate_subscriptions), (user_credit_balances), (credit_ledger)
Bookings	(bookings), (waitlist)
Feedback	(booking_feedback), (reviews)
Notifications	(notifications), (notification_templates)
AI/ML & Pinecone	(aiSuggestions), (nl_search_logs), (pinecone_sync_status), (chat_conversations), (chat_messages)

Domain	Tables
Compliance	<code>audit_logs</code> , <code>parental_consent_records</code> , <code>data_export_requests</code>
Platform	<code>platform_settings</code> , <code>invoices</code> , <code>feature_flags</code> , <code>api_keys</code>

6. JSONB Column Summary

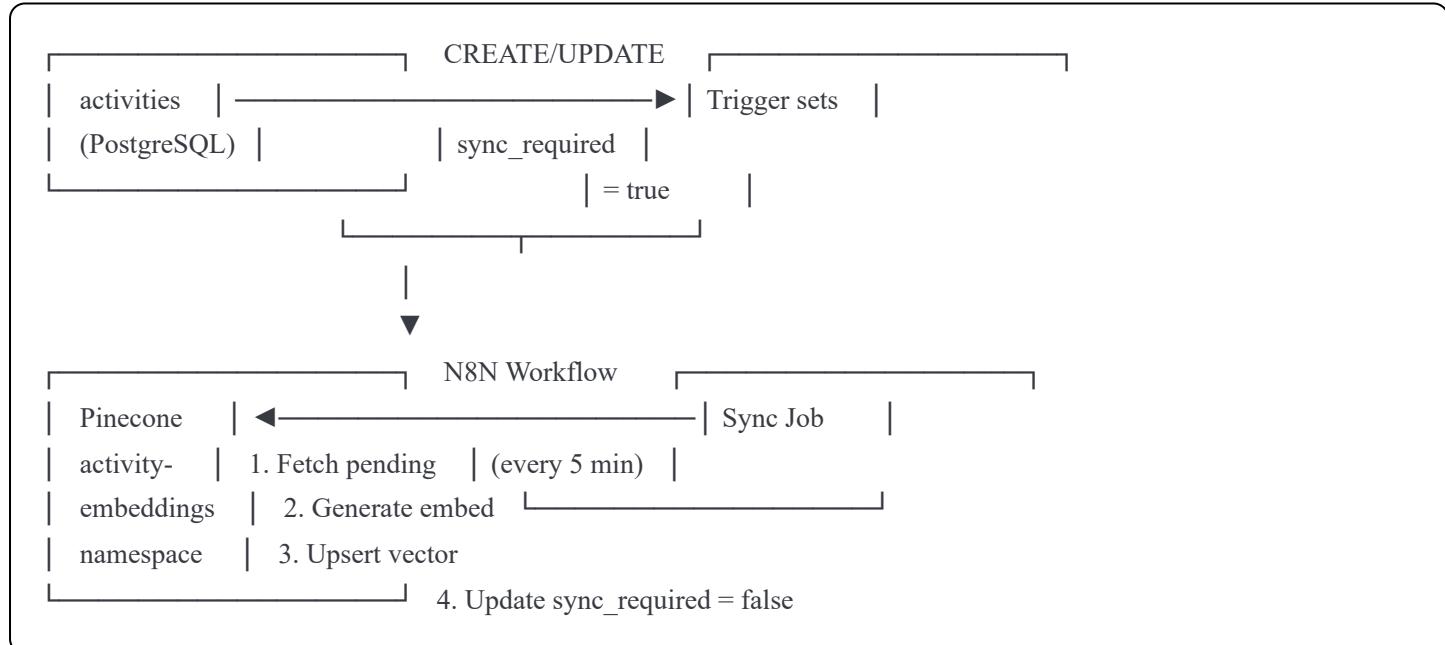
The following columns use JSONB for flexible, schema-less data storage:

Table	Column	Purpose
users	notification_preferences	User notification channel preferences
venues	business_hours	Operating hours by day of week
venues	gallery_images	Array of image URLs
venues	accessibility_features	Array of accessibility tags
activities	learning_outcomes	Array of learning outcome strings
activities	what_to_bring	Array of items to bring
activities	activity_tags	Key-value activity characteristics (synced to Pinecone metadata)
activities	accessibility_features	Array of accessibility tags
activity_recurring_schedules	days_of_week	Array of weekday numbers for schedule
children	interests	Array of interest categories (used in Pinecone embedding)
children	activities_to_avoid	Array of activity types to exclude (Pinecone filter)
children	accessibility_needs	Array of accessibility requirements
venue_performance_scores	improvement_suggestions	AI-generated improvement tips
venue_integrations	api_config	Integration-specific API configuration
aiSuggestions	child_ids	Array of child UUIDs for group suggestions
aiSuggestions	individual_fit_scores	Per-child Pinecone similarity scores
nl_search_logs	parsed_intent	LLM-extracted intent attributes
nl_search_logs	pinecone_filters	Filters applied to Pinecone query
nl_search_logs	result_activity_ids	Array of result UUIDs from Pinecone
nl_search_logs	result_scores	Pinecone similarity scores per result
chat_messages	parsed_intent	LLM-extracted intent for message
chat_messages	result_activity_ids	Activity results shown in chat
chat_messages	suggested_actions	Quick action chips shown to user

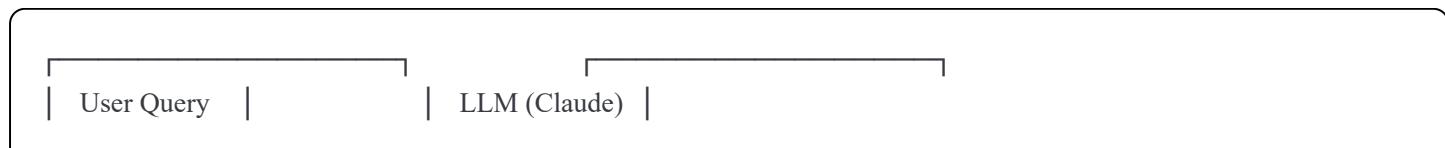
Table	Column	Purpose
notifications	data_payload	Rich notification context
notification_templates	available_variables	Merge variables for template
feature_flags	target_roles	Roles targeted by flag
feature_flags	target_corporate_ids	Corporates targeted by flag
feature_flags	target_user_ids	Users targeted by flag
feature_flags	variants	A/B test variant percentages
api_keys	scopes	Permission scopes for API key
api_keys	allowed_ips	IP whitelist for API key
audit_logs	old_values, new_values, context	Change tracking data
invoices	line_items	Invoice line item details

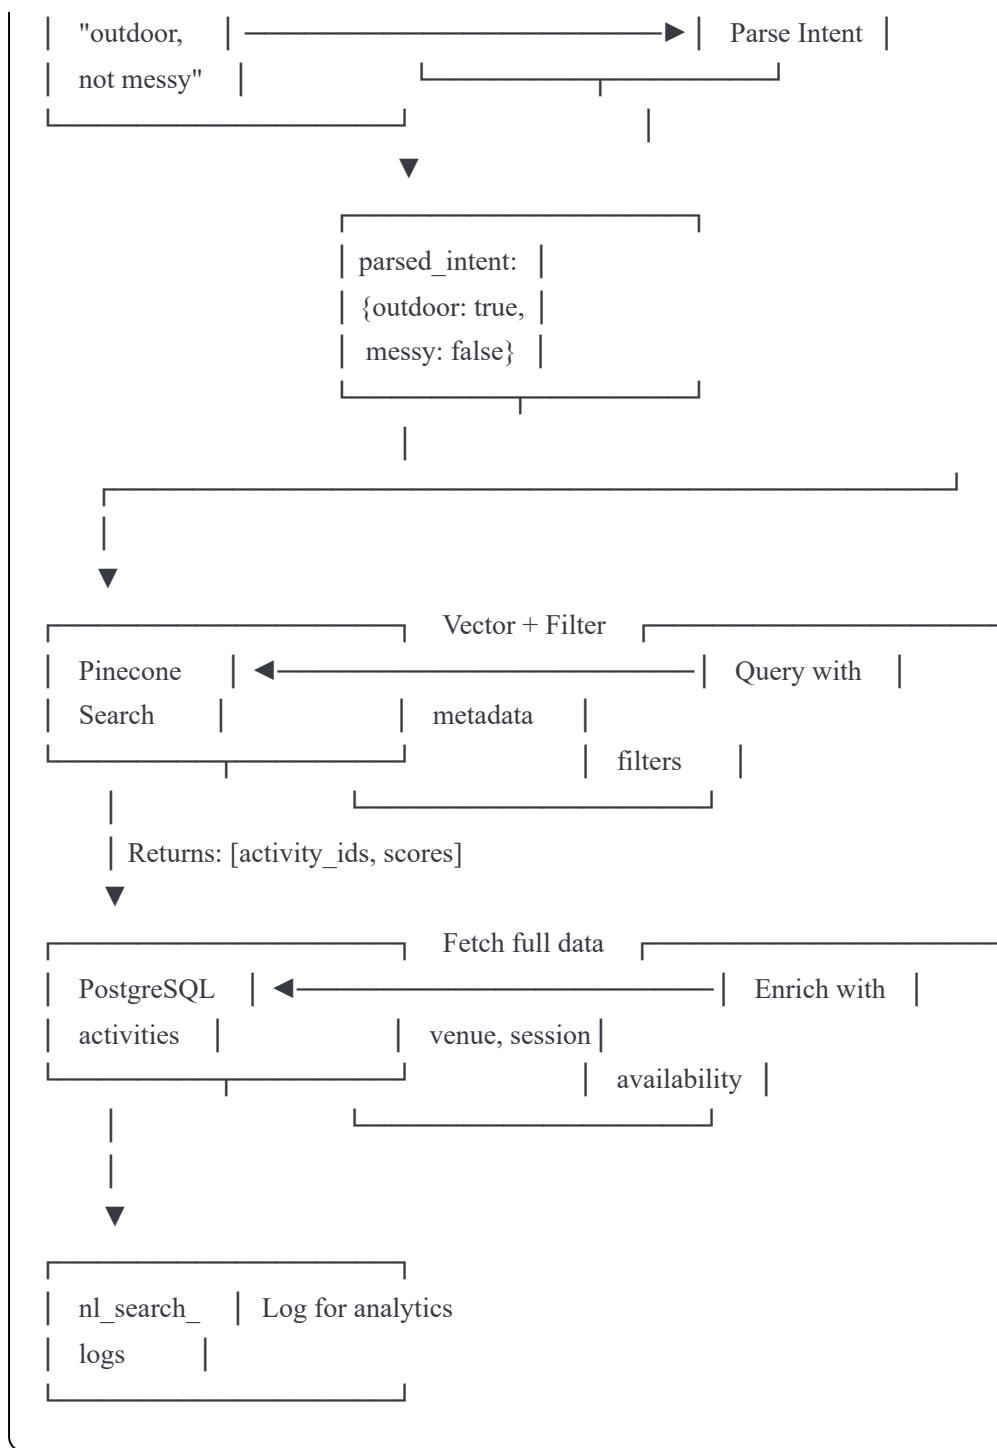
7. Pinecone ↔ PostgreSQL Data Flow

7.1 Activity Embedding Flow

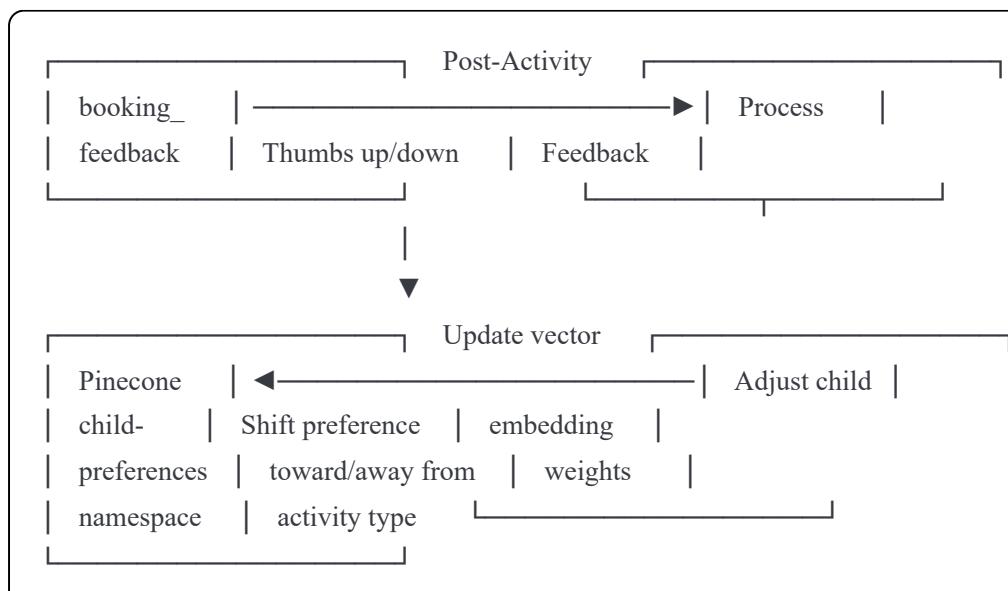


7.2 Search Query Flow





7.3 Child Preference Update Flow (Feedback Loop)



Update child	
feedback_	
incorporated	
_count	

Schema designed for Nexus Family Pass v7.0 - January 2026